

# OpenStack Operator's Manual

version 1

**Nick West**

April 22, 2021



# Contents

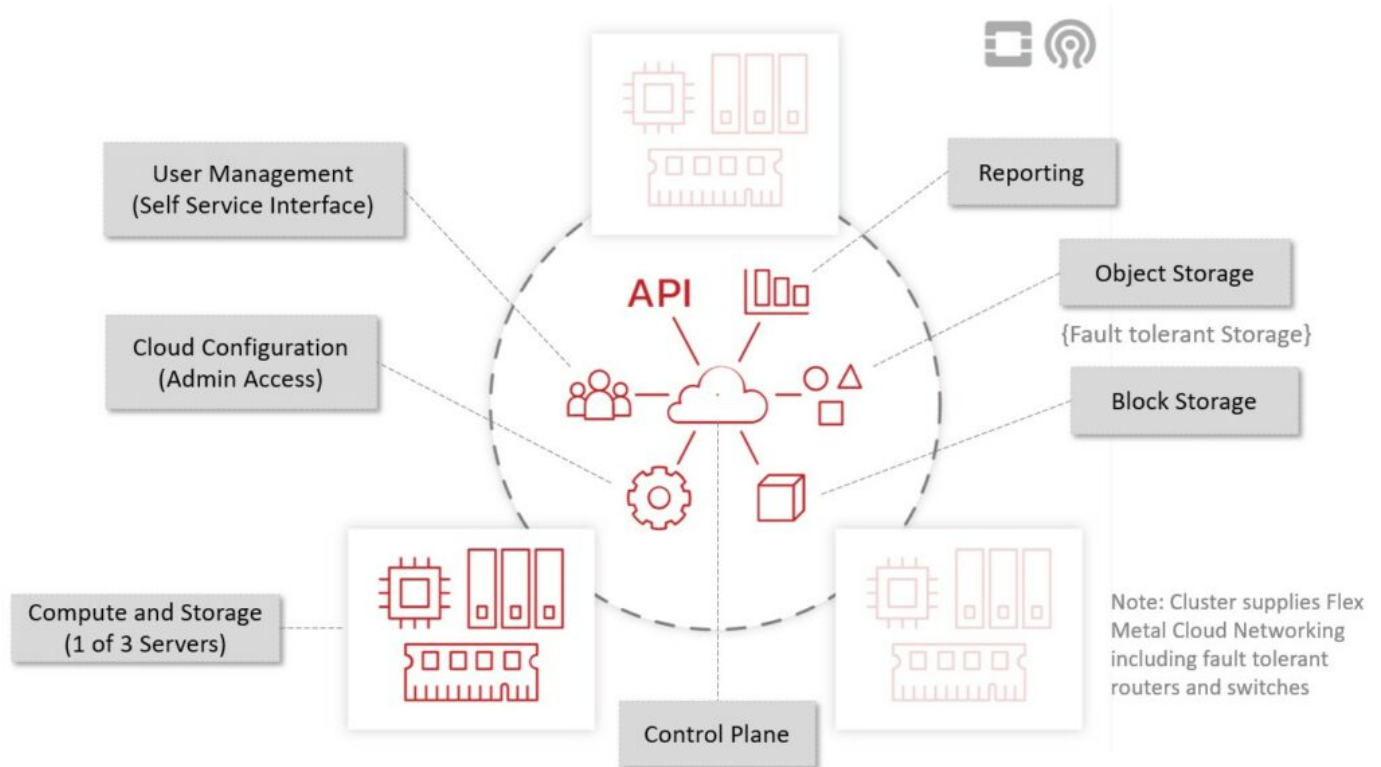
<b>OpenStack Operators Manual - Production in 40 Hours</b>	<b>1</b>
Private Cloud Core - OpenStack Operator Manual V1.2.1	1
<b>Getting Started with OpenStack</b>	<b>1</b>
Login to Horizon	1
Next Steps	3
<b>Install OpenStackClient</b>	<b>3</b>
Installing OpenStackClient	4
Next Steps	4
<b>Manage OpenStack Users, Groups and Projects</b>	<b>5</b>
Create and Manage Users	5
Create Users Using Horizon	5
Create Users Using the Command Line	6
Assign Role to a User	7
Create and Manage Groups	8
Create Groups Using Horizon	8
Add Users to a Group Using Horizon	8
Create Groups Using the Command Line	9
Add Users to a Group Using the Command Line	9
Create and Manage Projects	10
Create Projects Using Horizon	10
Create Projects Using the Command Line	11
Add Group to Project Using the Command Line	12
Next Steps	12
<b>Networks, Routers, and Floating IPs</b>	<b>12</b>
Exercises covered in this guide	12
Common terms	13
Create a network and router	13
Create a network	13
Create a router	14
Floating IPs	16
Allocate and Assign Floating IPs using Horizon	16
Troubleshooting	16
Check neutron docker containers	17
Next Steps	17
<b>Create a Network using OpenStackClient</b>	<b>17</b>
Exercises covered in this guide	17
Common terms	17
Create a network and router	18
Create a network	18
Create a router	19

Floating IPs	20
Allocate and Assign Floating IPs using OpenStackClient	20
Associate an instance with a private network	21
Troubleshooting	23
Check neutron docker containers	23
Next Steps	23
Create an OpenStack Instance using Horizon	23
Prerequisites for Creating an Instance	23
SSH key pairs	24
Create and upload an SSH key pair	24
Security groups	25
Create and manage security groups	25
Assign Storage	26
Create and attach a volume	26
Create an instance	26
Steps to create an instance	27
Troubleshooting Instance Error Status	29
Location of OpenStack service logs	29
Next Steps	29
Create an OpenStack Instance with OpenStackClient	29
OpenStackClient	30
Prerequisites for Creating an Instance	30
SSH key pairs	30
Create an SSH key pair	30
Security groups	31
Create and manage security groups	31
Assign Storage	32
Create and attach a volume	32
Create an instance	32
Before creating an instance	32
Procedure to create an instance	34
Troubleshooting	35
Troubleshooting Instance Error Status	35
Location of OpenStack service logs	35
Next Steps	35
SSH into an Instance	36
Access Instance associated with a Public IP	36
Using SSH	36
Using OpenStackClient	36
Access Instance associated with a Private IP	36
Next Steps	37
OpenStack Images	38

List and upload images	38
Using Horizon	38
List images	38
Upload and create images	38
Using the command line	39
List images	39
Upload an image	40
Create Image Snapshot	41
Snapshot using Horizon	41
Snapshot Using the command line	41
Create and verify an instance snapshot	41
Ceph, Cinder and Swift	42
Ceph	42
Ceph Data Durability	43
Cinder	43
More Information	43
What is block storage?	43
Swift	43
More Information	43
What is object storage?	43
How do Cinder and Swift relate to Ceph?	44
Using the Ceph CLI for Common Operational Tasks	44
Common tasks	44
More detailed ``ceph status`` output	44
Ceph Replication, Compression, and Erasure Coding	45
Ceph Replication	45
Setup replication across 2 OSDs	46
Setup replication across 3 OSDs	46
Erasure Coding	46
OpenStack CLI for Common Operational Tasks	46
Background	46
Get started	46
Output formatting	46
Common Tasks	47
Manage OpenStack users	47
Instance Management	48
Create an instance	48
Stop and start an instance	48
Create an instance snapshot	48
Perform live migration of instances	48
Troubleshoot instance issues	50
Upload images	50

Create a private network	51
Create security groups	51
Add SSH public key	52
Collect details about OpenStack environment	52
Backup OpenStack Service and Ceph Configurations	52
Background Information	52
Backup OpenStack services using kolla-ansible	53
Using kolla-ansible	53
Prepare and use kolla-ansible	53
Relevant files	53
Preparation	54
Backup kolla-ansible Files	54
Create OpenStack Database Backups	54
Restore OpenStack Database Backups	55
OpenStack Configuration Backups	56
Ceph Configuration Backups	56
Backing up Client Data and Disaster Recovery	56
Where is client data currently stored?	57
How to create backups of volumes	57
Using Horizon	57
Using OpenStackClient	58
Backing up ceph pool data	59
What to do in the event of a hard drive failure?	59
OpenStack Hardware Failures	59
What should be done should a hardware node fail?	60
Planned maintenance	60
Procedure for removing a compute node	60
Unplanned maintenance	61
Ceph failure scenarios and recovery	61
How do you know if a hard drive has failed?	61
Adding and Removing OpenStack Hardware Nodes	62
Types of Hardware Nodes	62
Adding a node	62
Removing a Storage and Compute Node	63
Requirements before removing a node	63
Remove a Storage and Compute node	63
Procedure	64
Automating OpenStack	69
Miscellaneous Guides	69
Associate an instance with a private network	70

# OpenStack Operators Manual - Production in 40 Hours



## Private Cloud Core - OpenStack Operator Manual V1.2.1

There are several phases of creating an OpenStack cloud and most of the public documentation focuses on the initial creation of a cloud. With Flex Metal, we provide a full private OpenStack so you will start as an "Operator". This guide, with a few noted exceptions, applies to any OpenStack that has been provisioned to OpenStack.org's RefStack standard.

This guide assumes you have followed the Flex Metal Central Process to provision your cloud already. Visit [Flex Metal Central to signup or sign in](#).

For product details please visit the [Flex Metal IaaS](#) or [Private Cloud](#) pages.

- [Start Here](#)

## Getting Started with OpenStack

To get started with OpenStack, you will first login to Horizon.

Horizon is a web-based interface used to administer OpenStack services, also known as OpenStack's dashboard. This is where you will do everything in OpenStack from making an instance, creating networks, managing your users, and so on.

For a full list of the purpose and features of Horizon, see OpenStack's [Horizon documentation](#).

## Login to Horizon

### Step 1 – Login to Flex Metal Central

To login to Horizon, you will first need to be logged into [Flex Metal Central](#).


**Step 2 – Navigate to the cloud**

Once logged in, click on the cloud you are working with to navigate to that cloud's details.

**Step 3 – Access Horizon**

Next, find and click the **Horizon** link on the left sidebar. It will be the last item in the list.

**Horizon link:**



Clicking this link will load the Horizon login page.

**Step 4 – How to obtain the Horizon password**

Initially, you will login to Horizon with the administrator account. The administrator account's username is **admin**.

Next, you will need to obtain the password for the **admin** account. The password is stored in a text file called `/etc/kolla/passwords.yml` on each hardware node.

To acquire the password use SSH to access a hardware node, then within `/etc/kolla/passwords.yml`, find the entry labeled **keystone\_admin\_password**. The file stores information in a “key: value” format, so the password will be the second entry on that line.

**Step 5 – SSH into a hardware node**

During cloud setup, you uploaded your SSH public key. This key exists on each hardware node and allows you to login to each node using SSH.

To SSH in, ensure you have one node's IP address, and the private key on the system you will use to access your cloud.

**Requirements to SSH into a node:**

- Username, which is root
- Authentication is with SSH keys
- An IP address of a hardware node

Here's an example command used to SSH in assuming `~/.ssh/YOUR_KEY` is the path to your private key and a hardware node IP is **50.50.50.50**:

```
$ ssh -i ~/.ssh/YOUR_KEY root@50.50.50.50
```

**Step 6 – Obtain Horizon password**

Once logged in, search for `keystone_admin_password` inside of `/etc/kolla/passwords.yml`.

For example:

```
# grep keystone_admin_password /etc/kolla/passwords.yml
keystone_admin_password: xxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxx
```



**Step 7 – Access Horizon**

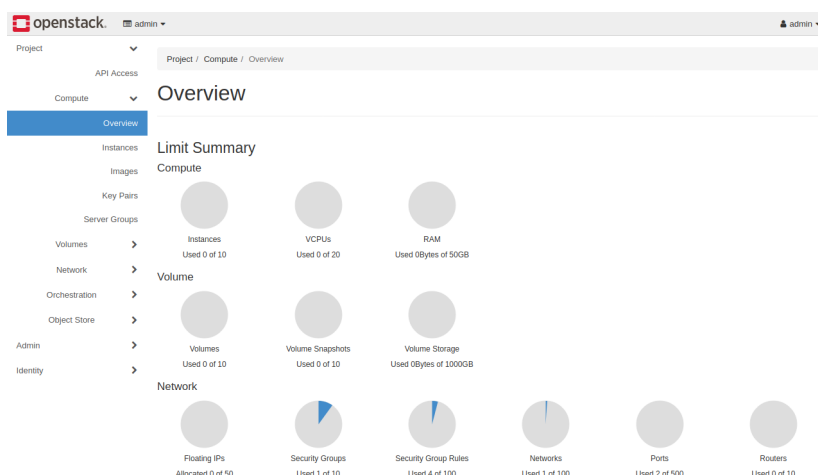
You can now login to Horizon using the credentials obtained from the previous sections.

**Horizon Login:**


The image shows the OpenStack Horizon login page. At the top is the OpenStack logo. Below it is the text "Log in". There are two input fields: "User Name" and "Password". The "Password" field has a toggle icon for visibility. At the bottom right is a blue "Sign In" button.

Enter in the credentials into the Horizon login page to login.

Once logged in, you should see something similar to:

**Initial Horizon login:****Next Steps**

The next guide in this series will cover installing OpenStack's command line client, OpenStackClient.

Should using the command line not be necessary for you, skip ahead to the guide that explains how to create Projects and Users.

**Install OpenStackClient**

OpenStack offers a way to administer your cloud over the command line using OpenStackClient.

If you would like to use OpenStackClient and are not familiar with how to install it, follow along with this installation guide.

**Note!** – These instructions are meant for a Linux environment.

## Installing OpenStackClient


OpenStackClient can be installed to any shell using pip. It can be installed locally to your machine or to an instance in the OpenStack cloud. You could think of the instance that has OpenStackClient installed as a “jumpstation”.

For more information on installing using pip, navigate to OpenStackClient's [pip project page](#).

It is encouraged to install OpenStackClient using a Python [virtual environment](#).

Once OpenStackClient is installed, you will need a copy of the **OpenStack RC file** as well as the OpenStack **clouds.yaml** file so API requests can be sent to your cloud. Both files can be found in Horizon under the **Project** tab, then **API Access**.

The section to download these files will appear like so:



Click the dropdown near the top right labeled **Download OpenStack RC File** and download the two files listed.

Place the OpenStack RC file in a safe place and `source` it using a shell:

```
$ source openrc.sh
```

Performing this step sets the needed environment variables in your shell so the API can be communicated with.

You will need to enter in your Horizon user's password upon sourcing this file.

The **clouds.yaml** file should be placed in `~/.config/openstack/clouds.yaml`.

Here is important information that explains where to place this file pulled from the top of an example **clouds.yaml** file:

```
# This is a clouds.yaml file, which can be used by OpenStack tools as a source
# of configuration on how to connect to a cloud. If this is your only cloud,
# just put this file in ~/.config/openstack/clouds.yaml and tools like
# python-openstackclient will just work with no further config. (You will need
# to add your password to the auth section)
# If you have more than one cloud account, add the cloud entry to the clouds
# section of your existing file and you can refer to them by name with
# OS_CLOUD=openstack or --os-cloud=openstack
```

You should have a working OSC to use after these steps.

Here's example usage of OSC where images are listed:

```
$ openstack image list
+-----+-----+-----+
| ID | Name | Status |
+-----+-----+-----+
| 6986aaf9-e602-4307-9c4f-6377795ff890 | CentOS 7 (ce7-x86_64) | active |
| d3f027f2-730f-4794-9288-44ed5c69050d | CentOS 8 (ce8-x86_64) | active |
+-----+-----+-----+
```

## Next Steps

Navigate to the next guide in this series for a walkthrough of how to manage OpenStack projects and users.

## Manage OpenStack Users, Groups and Projects

In OpenStack, managing your users, groups and projects is an important part of setting up the cloud. Projects are an organizational concept that allows a cloud to be subdivided. Users can be created and associated with a single or multiple projects. Groups are sets of users that can be assigned to projects.

Before getting into creating OpenStack components such as networks, instances, and the like, it is suggested you make use of projects so the cloud is organized as needed. In addition, you can create new user accounts and assign them to particular projects so your userbase has the appropriate access to the cloud.

This guide will explain how to create projects, additional user accounts, and groups.

### Create and Manage Users

When you start using OpenStack there is only the administrator user. You can think of this as the user “root” in a Linux environment. It has full privileges to the system. Due to this user having full privileges and that there is a chance to cause harm to a system with this user, it is suggested additional users be created as needed. The administrator user should typically be used only for tasks where that level of access is needed.

### Create Users Using Horizon

New OpenStack users can be created using Horizon. To create a new user, first ensure you are logged in as an administrator user. Typically, this is the user **admin** if you are just getting started.

Once logged in to Horizon, look for the section called **Identity**, then locate and click the **Users** link under that. This will bring up the page where users can be created.

#### OpenStack Users:

User Name	Description	Email	User ID	Enabled	Domain Name	Actions
admin	-		af82ee40927c4b72ad3011e77ab039e	Yes	Default	Edit
glance	-		05697a00f2242d39890621f33e81fb	Yes	Default	Edit

To make a new user, first click the **Create User** button near the top right to bring up the user creation form.

At minimum, you will need to enter a **User Name** and **Password** to create the user. In addition, an **Email** should be set as well.

All users have to be associated with a project, so choose the **Primary Project** this user will be assigned to.

Finally, select an appropriate **Role** for this user. Typically, the role of **member** is sufficient, however you can also assign the **admin** role. The **admin** role gives a user administrator access. In addition, custom roles can be created and assigned.

#### Create a User form:

With the user created, you should see it in the listing of users now.

#### User listing:

<input type="checkbox"/>	demo_user	-	demo@example.com	b7e2423b016b4debe5909f1b23f468	Yes	Default	Edit
--------------------------	-----------	---	------------------	--------------------------------	-----	---------	------

### Create Users Using the Command Line

The base command to create a user using OpenStackClient is:

```
$ openstack user create
```

Generally, when making a user using OpenStackClient, you will need to know the username, email address, and project to assign the user to.

Use `$ openstack project list` to list the project IDs.

#### List projects:

```
$ openstack project list
```

ID	Name
0d55c1cd820d4a5d9424456e1384ab73	Engineering
6a654535b8f04445bbc4974b2e4802cd	service
80eb7814893a414296ec1464d4a753b1	b9e8639372014c0b85cbfaffa6e1b5a8-a66df7d2-6e70-493f-922
b9e8639372014c0b85cbfaffa6e1b5a8	admin
c4006f982a2c4f63a2fabeeed6bc9f16	Project 1

This example will create a user called **demo\_user\_cli** and associate its default project to **Project 1**.

**NOTE!** – Entering passwords over the command line is generally considered insecure. You can pass the flag `--password-prompt` to interactively enter in the password.

#### Procedure:

The following demonstrates creating the **demo\_user\_cli** user:

```
$ openstack user create --project c4006f982a2c4f63a2fabeeed6bc9f16 \
--email demo@example.com --password-prompt demo_user_cli
User Password:
Repeat User Password:
```

Field	Value
default_project_id	c4006f982a2c4f63a2fabeeed6bc9f16
domain_id	default
email	demo@example.com
enabled	True
id	d88a89208d344cb4930761dd55a194d1
name	demo_user_cli
options	{}
password_expires_at	None

**List OpenStack users:**

```
$ openstack user list
```

ID	Name
af82ee40927c4b72ad3011e7fab03f9e	admin
05697a00ff2242d39890621f33e81fbb	glance
30c2e20a7c1141dc9fda9d405f1d6db3	cinder
ec64a60b1c6a4b64a559597417dd3ae2	placement
70b677b5fa4b4b8ca55699c8670f7993	nova
508ef9606a3a4048a86ec48e542020b4	neutron
5fce77bfae1a440d872d96982715af9e	heat
2da9eb178ee140c7aad2016f8d23ca9e	heat_domain_admin
e37f4e048f5e44c69305b6cec9ef2165	watcher
012b00425e9e4db289f2d71f6441d835	swift
b7e2423b016b4defbe5f09ff1b23f468	demo_user
d88a89208d344cb4930761dd55a194d1	demo_user_cli

From the above output, the **demo\_user\_cli** user is listed now.

**Assign Role to a User**

**NOTE!** – This section still needs to be filled out completely.

In OpenStack, there are roles that can be assigned to users and groups.

To view the current roles, use:

```
$ openstack role list
```

ID	Name
3e5d1f2c2c014bf0bfa929b0e31eb2b1	heat_stack_owner
764cb7fcf8214515860c628fcfb855d2	admin
aa2689853f7b42038afcdb797b54ef11	heat_stack_user
be24b48c21554b75849c66b9b710df84	reader
ccc32facd900440bb01a046db5c4096d	member
f5c0b887144d462bbd3bc35e9a0a9309	_member_

## Create and Manage Groups

Groups in OpenStack are collections of Users. These can be assigned to projects and make it easier to assign a grouping of users.

### Create Groups Using Horizon

To make a group, you will need to start with logging into to Horizon with an administrator account.

From there, on the left, find the **Identity** tab, then find the **Groups** tab within that. Following this will take you to the section where groups in OpenStack can be managed.

The following is a screenshot of how the **Groups** page will appear.

#### Group listing:

Groups

Group Name  Filter [+ Create Group](#) [Delete Groups](#)

Displaying 3 items

<input type="checkbox"/>	Name	Description	Group ID	Actions
<input type="checkbox"/>	Demo group		0f84a548d8cc42f7aeb5ba5d5f72ebe7	<a href="#">Manage Members</a>
<input type="checkbox"/>	Managed Hosting	Managed Hosting Team members	7085430cd7734bae8c54e38479300f0	<a href="#">Manage Members</a>
<input type="checkbox"/>	Development	Development team members	f77cec4a5aad4453ad28b8ba6562744	<a href="#">Manage Members</a>

Displaying 3 items

Next, find the **Create Group** button near the top right. This will load the form needed to create a group.

#### Create group:

Create Group

Name \*

Development

Description

Development team members

Description:

Groups are used to manage access and assign roles to multiple users at once. After creating the group, edit the group to add users.

Cancel

Create Group

Fill in the **Name** of the group and, optionally, a **Description**. Once the form has been submitted, the group will be created.

### Add Users to a Group Using Horizon

Now that a group has been made, users can be added to it.

To add users, pull up the listing of Groups in OpenStack first.

From there, find the group you are working with, and click **Manage Members** in the far right **Actions** column.

#### Manage Members:

<input type="checkbox"/>	Development	Development team members	f77cec4a5aad4453ad28b8ba6562744	<a href="#">Manage Members</a>
--------------------------	-------------	--------------------------	---------------------------------	--------------------------------

Clicking **Manage Members** will pull up another page.

#### Manage Members, page 2:

Identity / Groups / Group Management: Development

Group Management: Development


Filter  [+ Add Users](#)

User Name	Email	User ID	Enabled
No items to display.			

To add users, click the **Add Users** button near the top right.

A new form will appear providing a list of users that can be added.

#### Manage Members -> Add Users Form:



## Create Groups Using the Command Line

The base command to create a group using OpenStackClient is:

```
$ openstack group create
```

This section will demonstrate creating a group called **demo\_group**.

### Procedure

Use `$ openstack group create demo_group` to create the group:

```
$ openstack group create demo_group
+-----+-----+
| Field          | Value                                     |
+-----+-----+
| description    |                                           |
| domain_id     | default                                 |
| id            | 5d1177ede33b4cddadab6579408da7d7      |
| name          | demo_group                             |
+-----+-----+
```

### List groups:

```
$ openstack group list
+-----+-----+
| ID                                           | Name                               |
+-----+-----+
| 5d1177ede33b4cddadab6579408da7d7          | demo_group                         |
| 7085430cdf734bae8c54e384f79300f0          | Managed Hosting                    |
| f77cec4a5aad4453ad28b8fba6562744          | Development                        |
+-----+-----+
```

## Add Users to a Group Using the Command Line

Users can be added to groups using OpenStackClient. This will show an example where the user **demo\_user\_cli** is added to the group **demo\_group**.

The base command to add a user to a group is:

```
$ openstack group add user
```

### Procedure

Add the user **demo\_user\_cli** to the group **demo\_group**:

```
$ openstack group add user demo_group demo_user
```

Confirm the user was added successfully using `openstack group contains user`:

```
$ openstack group contains user demo_group demo_user
demo_user in group demo_group
```

## Create and Manage Projects

As an OpenStack administrator, it is typically advised that projects be created for specific uses. For example, you may want a project for development purposes, or need one for a specific department in your organization.

**NOTE!** – Projects can only be created by OpenStack accounts with **administrator** access.

This section will demonstrate how to create and manage projects using Horizon and the command line.

### Create Projects Using Horizon

To create a project, you will want to start with being logged into Horizon as an administrator account.

On the left, look for the section called **Identity**, then click on the **Projects** link under that. This page is where projects in OpenStack are managed.

To make a new project, say for the Engineering team in this example, find and click the **Create Project** button near the top right.

#### Create Project:

Name	Description	Project ID	Domain Name	Enabled	Actions
1b9e8639372014c0b85c0bffa1e1b5a8-e60464c5-9733-4c10-9e1f-a357cc9	Heat stack user project	14480a8ab224702b8dfce1ce09039a7	-	Yes	Manage Members
1b9e8639372014c0b85c0bffa1e1b5a8-e39564e0-baaf-496d-a73f-235aa12b	Heat stack user project	2a9fae1bad5a447bf7332113b84223	-	Yes	Manage Members

A form will display where the project's details will be needed.

#### Create Project Form:

On the first tab, enter a **Name** and a **Description** for the project.

The second tab, **Project Members**, allows you to add and remove members to this project. Here is where you will add the needed users to this project.

#### Create Project Form, page 2:

Finally, the third tab has to do with adding groups to this project. If you have a group to add, this is where it will be done.

#### Create Project Form, page 3:



Once all the details are filled in, click the **Create Project** button to create the project.

You should see the newly created project listed under the **Projects** page.

### Engineering project created:

<input type="checkbox"/> Name	Description	Project ID	Domain Name	Enabled	Actions
<input type="checkbox"/> Engineering	This is the engineering department's OpenStack project.	0d55c1cd82064a5d9424456e1384ab73	Default	Yes	Manage Members

## Create Projects Using the Command Line

The base command to create a project using OpenStackClient is:

```
$ openstack project create
```

This section will details the steps needed to create a project called **demo\_project**.

### Procedure

Create a project called **demo\_project**:

```
$ openstack project create demo_project
```

Field	Value
description	
domain_id	default
enabled	True
id	a7873c1cbbbe14607b5c5e797ef8d56ba
is_domain	False
name	demo_project
options	{}
parent_id	default
tags	[]

Confirm the project was created successfully using `openstack project show demo_project`:

```
$ openstack project show demo_project
```

Field	Value
description	
domain_id	default
enabled	True
id	a7873c1cbbbe14607b5c5e797ef8d56ba
is_domain	False
name	demo_project
options	{}
parent_id	default
tags	[]

## Add Group to Project Using the Command Line

Now that a group and project have been made, the group can be added to the project.

This section will demonstrate adding the group **demo\_group** to the project called **demo\_project**.

The base command to add a group to a project is:

```
$ openstack role add
```

### Procedure

Add group **demo\_group** to the project **demo\_project**:

```
$ openstack role add --project demo_project --group demo_group \
f5c0b887144d462bbd3bc35e9a0a9309
```

Verify the group was added to the project using `openstack role assignment list`:

```
$ openstack role assignment list --group demo_group --project demo_project --names
```

Role	User	Group	Project	Domain	System	Inherited
_member_		demo_group@Default	demo_project@Default			False

## Next Steps

With this topic covered, the next guide will explain how to create networks using Horizon in OpenStack.

Alternatively, if prefer working over the command line, see the how to create networks using OpenStackClient guide.

## Networks, Routers, and Floating IPs

This guide will explain basic networking functions in OpenStack including how to create a network, a router, and allocate and assign floating IPs.

The idea of this guide is it will explain how to create a private network with the intent to place virtual machines or instances on this network.

It is generally recommended that private networks be used where possible and to only expose the portions of your cloud to a public network when needed.

[Neutron](#) is the name of the service that handles networking in OpenStack. It provides “network connectivity as a service” between interfaces and uses the OpenStack Networking API.

Neutron allows networks, routers, floating IPs, and security groups to be created.

## Exercises covered in this guide

This guide will explain how to:

- Create a private network
- Create a router
- Associate an instance with a private network
- Allocate floating IPs
- Assign a floating IP to an instance

## Common terms

- **Provider Network**
  - a network that has been mapped to physical networking devices
  - this network comes already setup and is Internet-accessible
- **Floating IP**
  - public facing and allows external communication
  - attach to an instance on a private network to allow access to the Internet
  - allocated from the provider network
- **Port**
  - typically created as a result of another action (creating an instance)
  - associated with instances, routers, floating IPs, and essentially anything that can be connected to a network

## Create a network and router

Networks and routers can be created in OpenStack. To make a private network accessible from the provider network, a router must be created.

This section will go over how to make a network and router in Horizon, then connect the public network to a private network using the router.

## Create a network

To make a network in Horizon, find the **Network** tab on the left, then navigate to the **Networks** tab under that. Finally, locate the **Create Network** button near the top right. You'll be presented with the form to create a network.

### Create a Network:

Create Network

Network Subnet Subnet Details

Network Name  
Internal

Create a new network. In addition, a subnet associated with the network can be created in the following steps of this wizard.

☒ Enable Admin State ⓘ

☐ Shared

☒ Create Subnet

Availability Zone Hints ⓘ  
nova

Cancel Back Next >

Specify a name for the network under **Network Name** on the first page. In this example, the network will be called **Internal**. Fill out any other needed details and navigate to the **Subnet** tab.

### Create a Subnet:

**Create Network**

Network Subnet Subnet Details

Subnet Name:

Network Address:

IP Version:

Gateway IP:

☐ Disable Gateway

Cancel Back Next

Creates a subnet associated with the network. You need to enter a valid "Network Address" and "Gateway IP". If you did not enter the "Gateway IP", the first value of a network will be assigned by default. If you do not want gateway please check the "Disable Gateway" checkbox. Advanced configuration is available by clicking on the "Subnet Details" tab.

- Under **Subnet Name** specify a name for the subnet
- For **Network Address** choose a network in CIDR notation. This example uses **192.168.0.0/24**.
- Finally under **Gateway IP** specify the gateway IP for this network. If the gateway IP is not filled out, one will be chosen by the neutron service.

The final tab is called **Subnet Details**. This tab does not need to be filled out to create the network and subnet. This example will stop here, however on this page you can enable or disable DHCP, specify specific IPs to be allocated, set DNS name servers, and set Host Routes.

With the previous steps done, the network has been created. Loading the **Network -> Networks** tab will display the new network:

**Internal network created:**

Network  Filter [+ Create Network](#) [Delete Networks](#)

Network Topology ☐ Networks ☐ Routers

Displaying 2 items

<input type="checkbox"/>	Name	Subnets Associated	Shared	External	Status	Admin State	Availability Zones	Actions
<input type="checkbox"/>	Internal	internal_subnet 192.168.0.0/24	No	No	Active	UP	nova	<a href="#">Edit Network</a>

## Create a router

With a network created, the next step is to create a router which will bridge the connection from the **External or provider network** to the private network.

To make a router in Horizon, find the **Network** tab on the left, then locate the **Routers** tab under that. This page will list current routers and allows you to create a router.

To create a new router, click the **Create Router** button near the top right.

**Create a router:**

Routers

Filter [+ Create Router](#)

Name	Status	External Network	Admin State	Availability Zones	Actions
No items to display.					

**Create a router form:**

**Create Router**

Router Name:

☒ Enable Admin State

External Network:

☒ Enable SNAT

Availability Zone Hints:

Cancel Create Router

Description:  
Creates a router with specified parameters.  
Enable SNAT will only have an effect if an external network is set.

Under **Router Name** choose a name for the router. This example router will be called **router\_1**.

The router will need to be connected to an external network, which will be the provider network. The network called **External** will be used.

Once created it will show in the list of routers.

### Router listing:

Routers

Router Name  Filter [+ Create Router](#) [Delete Routers](#)

Displaying 1 item

<input type="checkbox"/>	Name	Status	External Network	Admin State	Availability Zones	Actions
<input type="checkbox"/>	router_1	Active	External	UP	nova	<a href="#">Clear Gateway</a>

This takes care of creating the router with the **External** network attached. To allow the router to communicate with the **External** and the **Internal** network, the **Internal** network will need to be attached. The remaining steps demonstrate how to do this.

### Attach **Internal** network

First, pull up the newly created router from the **Network -> Routers** section of Horizon.

Next, to add an interface to this router, click the router's name in the listing of routers, then click the **Add Interface** button near the top right.

### Add Interface to Router:

Add Interface

Subnet \*

Internal: 192.168.0.0/24 (internal\_subnet) ▼

IP Address (optional) ⓘ

Description:

You can connect a specified subnet to the router.

If you don't specify an IP address here, the gateway's IP address of the selected subnet will be used as the IP address of the newly created interface of the router. If the gateway's IP address is in use, you must use a different address which belongs to the selected subnet.

Cancel

Submit

In the form that appears choose the subnet to connect the router to and optionally choose an IP address for the interface. If you don't choose an IP, the gateway IP of the subnet will be used.

With the interface added to the router, the **External** and **Internal** networks are now connected. This can visually be seen by navigating to the **Network** tab on the left, then the **Network Topology** tab under that.

### Network Topology:




## Floating IPs

Floating IPs in OpenStack are publicly routable IP addresses that can be attached and detached to instances. For example if there's an instance associated with a private network but needs to be accessed from the Internet, a floating IP can be associated with the instance, allowing communication from the Internet.

### Allocate and Assign Floating IPs using Horizon

To use Floating IPs they will first need to be allocated from the provider network's pool of IPs.

To allocate floating IPs in Horizon, navigate to the **Network** tab on the left, and look for **Floating IPs**.



Click the **Allocate IP To Project** to allocate a new IP.

**Allocate IP To Project:**

The IP will be obtained from the provider network. In this example, the name of that network is **External**.

Once the IP is added, it will appear in the floating IP list.

**Floating IP list:**

<input type="checkbox"/>	IP Address	Description	DNS Name	DNS Domain	Mapped Fixed IP Address	Pool	Status	Actions
<input type="checkbox"/>	50.50.50.50	Web Server			-	External	Down	Associate

Displaying 1 item

Now that an IP has been allocated, it can be assigned to an instance. Use the **Associate** button from the drop down on the right.

**Associate floating IP:**

Choose the floating IP and the port to associate it to.

## Troubleshooting

**NOTE!** – This section will be need to be updated and is incomplete.

### Check neutron docker containers

In Private Cloud Core, OpenStack has been deployed using `kolla-ansible`. This means each service is running in a Docker container on the hardware nodes.

You can check the status of the neutron Docker containers to see if they are running:

```
# docker ps | grep -i neutron
66c57a96ecf5      kolla/centos-binary-neutron-metering-agent:train-centos8      "dumb-init
b407dda02505      kolla/centos-binary-neutron-metadata-agent:train-centos8      "dumb-init
93ce576891db      kolla/centos-binary-neutron-l3-agent:train-centos8             "dumb-init
5531d47d6405      kolla/centos-binary-neutron-dhcp-agent:train-centos8           "dumb-init
d3e20fd98aef      kolla/centos-binary-neutron-openvswitch-agent:train-centos8    "dumb-init
2802a09606e5      kolla/centos-binary-neutron-server:train-centos8              "dumb-init
```

### Next Steps

With this guide complete, move on to the next guide, which explains how to create an Instance in Horizon.

Alternatively, if you prefer working with the command line, see the guide that explains how to create an Instance using `OpenStackClient`.

### Create a Network using OpenStackClient

This guide will explain basic networking functions in OpenStack including how to create a private network, a router, and allocate and assign floating IPs.

[Neutron](#) is the name of the service that handles networking in OpenStack. It provides “network connectivity as a service” between interfaces and uses the OpenStack Networking API.

Neutron allows networks, routers, floating IPs, and security groups to be created.

### Exercises covered in this guide

This guide will explain how to:

- Create a private network
- Create a router
- Associate an instance with a private network network
- Allocate floating IPs
- Assign a floating IP to an instance

### Common terms

- **Provider Network**
  - a network that has been mapped to physical networking devices
  - this network comes already setup and is Internet-accessible

- Floating IP
    - public facing and allows external communication
    - attach to an instance on a private network to allow access to the Internet
    - allocated from the provider network
  - Port
    - typically created as a result of another action (creating an instance)
    - associated with instances, routers, floating IPs, and essentially anything that can be connected to a network
- 

### Create a network and router

Networks and routers can be created in OpenStack. To make a private network accessible from the provider network, a router must be created.

This will go over how to make a network and router using the command line with OpenStackClient.

---

### Create a network

Listed are the steps needed to create a private network. Variables are presented in all capital and should be replaced accordingly. Note the output of most of the commands has been truncated.

Use this command to create a network, replacing NETWORK\_NAME with the name of the network:

```
$ openstack network create NETWORK_NAME
```

Create a network called **private\_network**:

```
$ openstack network create private_network --fit-width
```

Field	Value
admin_state_up	UP
availability_zone_hints	
availability_zones	
created_at	2020-10-30T20:57:16Z
description	
dns_domain	
id	890e5ab6-a5d7-44c2-bf84-351342006cdd

Next, a subnet will need to be created.

Use this command to create a subnet, replacing NETWORK\_NAME and SUBNET\_NAME with the respective names of the network and subnet and replace SUBNET\_RANGE with the subnet to use. An example subnet range could be **10.0.0.0/24**:

```
$ openstack subnet create --subnet-range SUBNET_RANGE --network NETWORK_NAME \
SUBNET_NAME
```

Create a subnet called **private** with subnet range of **10.0.0.0/24**:

```
$ openstack subnet create --subnet-range 10.0.0.0/24 --network private_network private --fit-width
```

Field	Value
allocation_pools	10.0.0.2-10.0.0.254



cidr	10.0.0.0/24
created_at	2020-10-30T21:03:50Z
description	
dns_nameservers	
dns_publish_fixed_ip	None
enable_dhcp	True
gateway_ip	10.0.0.1
host_routes	
id	46d0b88c-f1cf-4e40-a395-9281a7dd59d9

### Create a router

The following are the commands required to create a router.

To make a router, use this base command, replacing `ROUTER_NAME` with the name of the router:

```
$ openstack router create ROUTER_NAME
```

Create a router called **router\_1**:

```
$ openstack router create router_1 --fit-width
```

Field	Value
admin_state_up	UP
availability_zone_hints	
availability_zones	
created_at	2020-10-30T21:08:38Z
description	
distributed	False
external_gateway_info	null
flavor_id	None
ha	True
id	a700e1b1-36ac-4e4f-b4c7-fa80be513656

With the router created, a subnet needs to be attached to it along with the external or, provider network.

To add a subnet, use this command, replacing `ROUTER_NAME` and `SUBNET_NAME` with the names of the respective router and subnet:

```
$ openstack router add subnet ROUTER_NAME SUBNET_NAME
```

Add subnet **private** to the router called **router\_1**:

```
$ openstack router add subnet router_1 private
```

The command to add the subnet to the router returns no output if successful.

Finally the router also needs the external network connected to it.

Use this command to add an external network, replacing `EXTERNAL_NETWORK_UUID` with the UUID of the network:

```
$ openstack router set --external-gateway EXTERNAL_NETWORK_UUID \
ROUTER_NAME
```

You can get the UUID of the network to use by running `openstack network list`. The UUID will be listed in the first column.

Add the external network to the router called **router\_1**:

```
$ openstack router set --external-gateway \
55d31bd5-77ba-4ed0-ab6e-99554b33aa90 router_1
```

With these steps completed, you have a router that connects the external network to the private network.

You can see the details of the router by running this, replacing `ROUTER_NAME` with the name of the router:

```
$ openstack router show ROUTER_NAME
```

Show the details for the router called **router\_1**, including the interfaces that were previously attached:

```
$ openstack router show router_1 --fit-width
```

Field	Value
admin_state_up	UP
availability_zone_hints	
availability_zones	nova
created_at	2020-10-30T21:08:38Z
description	
distributed	False
external_gateway_info	{ "network_id": "55d31bd5-77ba-4ed0-ab6e-99554b33aa90", "external_network_subcidr": "173.231.202.88" }, { "subnet_id": "ca55b96b-0d70-4cbc-92fc-7bf5ce74cfa4", "ip_address": "173.231.202.88" } ], "enable_snat": true }
flavor_id	None
ha	True
id	a700e1b1-36ac-4e4f-b4c7-fa80be513656
interfaces_info	[ { "port_id": "8c45b709-9f74-486d-bb13-0d8b459066dd", "ip_address": "169.254.169.169", "subnet_id": "52df9b51-8a5a-45f0-bf22-693eb9712c32" }, { "port_id": "d0233252-3756-44b2-8bad-e3e82781f96d", "ip_address": "169.254.169.169", "subnet_id": "52df9b51-8a5a-45f0-bf22-693eb9712c32" }, { "port_id": "e4bd63c8-8272-4285-85b2-f98e9f028471", "ip_address": "169.254.169.169", "subnet_id": "52df9b51-8a5a-45f0-bf22-693eb9712c32" }, { "port_id": "e8d5407c-1574-45e4-8f04-77dc1a848591", "ip_address": "10.0.0.1", "subnet_id": "46d0b88c-f1cf-4e40-a395-9281a7dd59d9" } ]

## Floating IPs

Floating IPs in OpenStack are publicly routable IP addresses that can be attached and detached to instances. For example if there's an instance associated with a private network but needs to be accessed from the Internet, a floating IP can be associated with the instance, allowing communication from the Internet.

## Allocate and Assign Floating IPs using OpenStackClient

OpenStackClient can be used to manage Floating IPs. The following is a list of commands used to manage floating IPs.

**Allocate additional floating IPs where `NETWORK` is the UUID of the network to allocate IPs from:**

```
$ openstack floating ip create NETWORK
```

```
$ openstack floating ip create 55d31bd5-77ba-4ed0-ab6e-99554b33aa90 --fit-width
```

Field	Value
created_at	2020-10-29T20:05:53Z
description	
dns_domain	

dns_name	
fixed_ip_address	None
floating_ip_address	50.50.50.50
floating_network_id	55d31bd5-77ba-4ed0-ab6e-99554b33aa90
id	e1818df0-ce58-4f36-94ff-170a5a0c1f9f

Some of the output is truncated.

View floating IPs:

```
$ openstack floating ip list --fit-width
```

ID	Floating IP Address	Fixed IP Address	Port	Floating Network
e1818df0-ce58-4f36-94ff-170a5a0c1f9f	50.50.50.50	None	None	55d31bd5-77ba-4ed0-ab6e-99554b33aa90

Assign floating IPs to an instance:

```
$ openstack floating ip set FLOATING_IP
```

### Associate an instance with a private network

This section will explain how to take an instance created on the provider network and associate that with a private network. This may not be a very common task assuming an instance was created on the appropriate network to begin with.

In the event an instance was created but on the wrong network, you can have it associated with another network.

The commands needed to add and remove networks an instance is associated with are:

- `openstack server add network SERVER NETWORK`
- `openstack server remove network SERVER NETWORK`

SERVER and NETWORK are placeholders and should be replaced with the UUIDs of the server and network needed.

To change the network for instance, some information is needed first:

- Instance UUID
- UUID of network to which the instance is moving
- UUID of network from which the instance is moving

To obtain the UUID of the instance use:

```
$ openstack server list
```

Obtain UUID of instance called **Server 1**:

```
$ openstack server list --fit-width
```

ID	Name	Status	Networks	Image
4208022c-3afa-4233-84ed-8df04bb2c4ce	Server 1	ACTIVE	External=173.231.202.87	

This instance is shown to be on the **External** network and instead it should be associated with another network, called **private\_network**.

To move this instance to the **private\_network**, the UUID of both networks is needed.

Use the following to list network UUIDs:

```
$ openstack network list
```

List the UUID of each network:

```
$ openstack network list --fit-width
```

ID	Name	Subnets
3eb74273-45aa-45e7-977e-3c706c0499fb	Internal	c6bb7ad7-79
55d31bd5-77ba-4ed0-ab6e-99554b33aa90	External	ca55b96b-0d
8346bc1d-2b69-421d-ab7f-008b9af53c5d	HA network tenant	52df9b51-8a
890e5ab6-a5d7-44c2-bf84-351342006cdd	private_network	46d0b88c-fl

From this output, the UUIDs can be obtained.

Next the instance will need to be associated with the network called **private\_network**.

Associate **Server 1** with **private\_network** using the UUIDs from above:

```
$ openstack server add network 4208022c-3afa-4233-84ed-8df04bb2c4ce \
890e5ab6-a5d7-44c2-bf84-351342006cdd
```

Disassociate **Server 1** from the **External** network:

```
$ openstack server remove network 4208022c-3afa-4233-84ed-8df04bb2c4ce \
55d31bd5-77ba-4ed0-ab6e-99554b33aa90
```

Both the add and remove network commands do not return any output if they are successful.

To be sure the network has been associated correctly, use `openstack server show SERVER_UUID` to confirm, replacing `SERVER_UUID` with the UUID of the server you are working with.

Confirm new network for **Server 1**:

```
$ openstack server show 4208022c-3afa-4233-84ed-8df04bb2c4ce --fit-width
```

Field	Value
OS-DCF:diskConfig	AUTO
OS-EXT-AZ:availability_zone	nova
OS-EXT-SRV-ATTR:host	hcl.example.com
OS-EXT-SRV-ATTR:hypervisor_hostname	hcl.example.com
OS-EXT-SRV-ATTR:instance_name	instance-00000096
OS-EXT-STS:power_state	Running
OS-EXT-STS:task_state	None
OS-EXT-STS:vm_state	active
OS-SRV-USG:launched_at	2020-10-23T20:52:15.000000
OS-SRV-USG:terminated_at	None
accessIPv4	
accessIPv6	
addresses	private_network=10.0.0.192
config_drive	
created	2020-10-23T20:51:27Z
flavor	hcl.small (hcl.small)
hostId	fc32f234767ef0316a30779802b178074140cb60635ff81fa1cd
id	4208022c-3afa-4233-84ed-8df04bb2c4ce

The **addresses** field confirms this instance is now associated with the network called **private\_network**.

## Troubleshooting

**NOTE!** – This section will be need to be updated and is incomplete.

---

### Check neutron docker containers

In Private Cloud Core, OpenStack has been deployed using `kolla-ansible`. This means each service is running in a Docker container on the hardware nodes.

You can check the status of the neutron Docker containers to see if they are running:

```
# docker ps | grep -i neutron
66c57a96ecf5      kolla/centos-binary-neutron-metering-agent:train-centos8      "dumb-init
b407dda02505      kolla/centos-binary-neutron-metadata-agent:train-centos8      "dumb-init
93ce576891db      kolla/centos-binary-neutron-l3-agent:train-centos8             "dumb-init
5531d47d6405      kolla/centos-binary-neutron-dhcp-agent:train-centos8           "dumb-init
d3e20fd98aef      kolla/centos-binary-neutron-openvswitch-agent:train-centos8    "dumb-init
2802a09606e5      kolla/centos-binary-neutron-server:train-centos8               "dumb-init
```

---

## Next Steps

The next guide will explain how to make an instance.

With this guide complete, you should have all that is needed to create an instance.

Navigate to the create an instance with the command line guide to continue this series.

Should you want to learn to the make the instance using Horizon, see the create an instance using Horizon guide.

## Create an OpenStack Instance using Horizon

The purpose of this guide is to explain how to create an instance using Horizon. An example will be demonstrated where an instance is created on a private network. The intent of the instance will be to act as a “jumpstation” where `OpenStackClient` will be installed. Finally, this jumpstation can be used to administer the OpenStack cloud using `OpenStackClient`.

You'll also learn how to upload or create an SSH key pair, assign storage using a volume, and create a security group. Each of these components will be added to the instance.

Make sure you are logged into Horizon when following this guide. If you are not familiar with accessing Horizon, see the Getting Started guide.

---

## Prerequisites for Creating an Instance

Before making an instance, the following should exist:

- A flavor
- An image
- A network
- A security group
- An SSH public key

At minimum, a flavor, an image, and a network should exist. You may also want to attach a volume, specify security groups to use, and attach an SSH public key.

## SSH key pairs


An SSH key pair will be required to access any instances over SSH. Password authentication is by default disabled in the operating system images.

You can either have an SSH key created or you can upload the public key of your SSH key pair.

### Create and upload an SSH key pair

To make an SSH key pair, start in the main dashboard of Horizon. On the left, click the **Compute** heading. Following that, find the **Key Pairs** link under **Compute**. Clicking this link takes you to the SSH Key Pairs section.

**SSH key pair section:**




To make an SSH key pair, click **Create Key Pair** near the top and to the right. You'll need to fill out the key pair name and the type of key it will be. Once done, you'll be prompted to download the private key which needs to be stored in a safe place. Only you should have access to your private SSH key pair! The public key portion of the key pair now exists in Horizon.

**Create SSH key pair form:**

Should you already have an SSH key pair you want to use, you can instead upload your public key. This can be done in the same **Key Pairs** section. Locate the **Import Key Pair** button near the top right of the screen. You'll need to fill out the name of the key pair, specify the type of key it is, and finally choose the public key from your computer or paste the public key in.

**Upload SSH key pair form:**

After an SSH key has been created or uploaded, they should appear in the list like so:



## Security groups

A security group in OpenStack controls network access. If no security group is assigned, a default one will be used. The default security group disables all inbound traffic and only allows outbound traffic, which may not be useful.

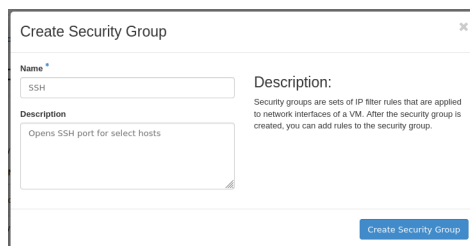
In Horizon, security groups can be found on the left, under **Network**, then **Security Groups**. You can create and modify security groups using this section.

A security group should be made that allows the type of incoming and outgoing traffic required. OpenStack's default security group does not allow incoming traffic so if you wanted to SSH into an instance, port 22 will be closed. An example of a security group would be one that allows web traffic by opening ports 80 and 443 and you may want another security group that opens port 22 for SSH traffic.

## Create and manage security groups

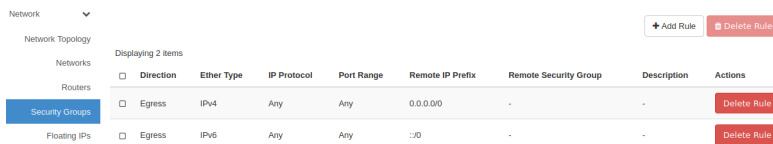
In Horizon, on the main landing page, look for the **Network** tab, then find **Security Groups**. Click [+ Create Security Group](#) to load the form to create a Security Group.

### Create Security Group form:



Fill out the name and description of the group. Once done, you'll be taken to a page where you can define egress (outbound) and ingress (inbound) rules for the group.

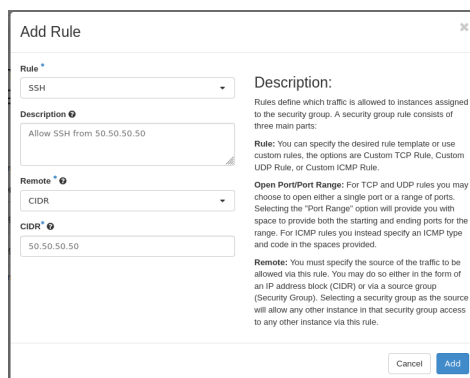
### Manage security group rules:



Direction	Ether Type	IP Protocol	Port Range	Remote IP Prefix	Remote Security Group	Description	Actions
Egress	IPv4	Any	Any	0.0.0.0	-	-	Delete Rule
Egress	IPv6	Any	Any	::0	-	-	Delete Rule

Let's say you want to allow incoming SSH traffic from 50.50.50.50. To do so, click the **Add Rule** button near the top right.

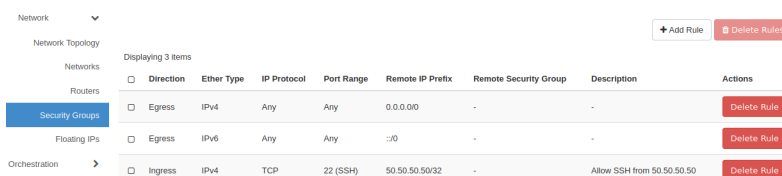
### Add Rule:



From here you can specify a custom rule using the various options, or you can use the first drop down to specify common protocols, such as SSH.

To add the rule for SSH, select from the first drop down the **SSH** option. You can describe the rule using the next box. The third field called **Remote** should be left as **CIDR**. Finally the last box needs to be filled in with the IP, 50.50.50.50. You can specify a range of IPs using CIDR notation as well.

#### Newly added SSH rule:



The screenshot shows the 'Security Groups' tab in the OpenStack dashboard. A table lists three security rules. The third rule is the newly added SSH rule.

Direction	Ether Type	IP Protocol	Port Range	Remote IP Prefix	Remote Security Group	Description	Actions
Egress	IPv4	Any	Any	0.0.0.0/0	-	-	Delete Rule
Egress	IPv6	Any	Any	:::0	-	-	Delete Rule
Ingress	IPv4	TCP	22 (SSH)	50.50.50.50/32	-	Allow SSH from 50.50.50.50	Delete Rule

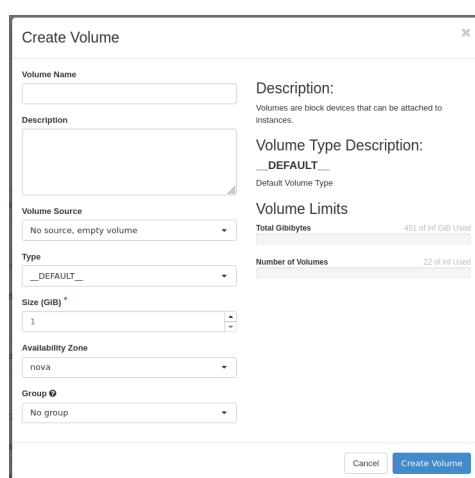
## Assign Storage

Storage can be assigned through volumes. A volume in OpenStack is like a removable USB drive that can be attached to instances as seen fit. Cinder is the OpenStack service that allows volumes to be created. A volume can also be used to boot an instance.

### Create and attach a volume

On the main page of Horizon, on the left, under the **Volumes** tab, look for the **Volumes** sub tab. From here you may see a list of volumes and can also create one. To create a volume, click the **+ Create Volume** icon.

#### Create volume form:



The 'Create Volume' form contains the following fields and sections:

- Volume Name:** Text input field.
- Description:** Text area.
- Volume Source:** Dropdown menu with 'No source, empty volume' selected.
- Type:** Dropdown menu with '\_DEFAULT\_' selected.
- Size (GiB):** Text input field with '1' entered.
- Availability Zone:** Dropdown menu with 'nova' selected.
- Group:** Dropdown menu with 'No group' selected.
- Description:** Text area with placeholder text: 'Volumes are block devices that can be attached to instances.'
- Volume Type Description:** Text area with placeholder text: 'Default Volume Type'.
- Volume Limits:**
  - Total GiB/bytes:** Progress bar showing 451 of inf GiB Used.
  - Number of Volumes:** Progress bar showing 22 of inf Used.
- Buttons:** 'Cancel' and 'Create Volume'.

Minimally, the size of the volume needs to be specified. All other options are optional.

There are several sources that can be used for volumes, including other volumes, images, and snapshots.

To create a new empty volume, specify "No source, empty volume" under the **Volume Source** option.

## Create an instance

An instance is another name for a virtual machine in OpenStack. Instances are created by the Nova service and contribute to the processing power of the cloud.

With the previous steps followed, you should have everything needed to make an instance.

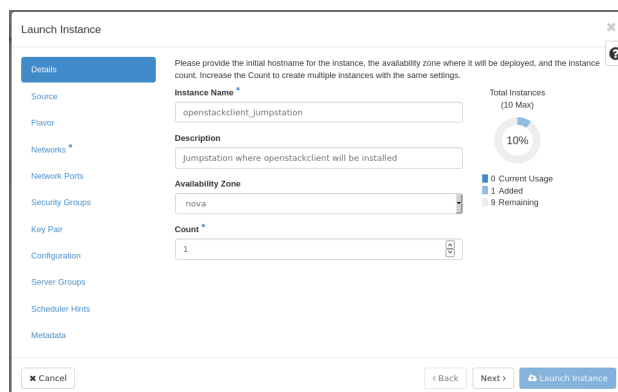


This section will demonstrate creating a “jumpstation” instance on a private network.

## Steps to create an instance

In Horizon, look for the **Compute** heading on the left. From there, choose **Instances**. This displays the Instances page. Click the  button to create one.

### Launch Instance



Launch Instance

Please provide the initial hostname for the instance, the availability zone where it will be deployed, and the instance count. Increase the Count to create multiple instances with the same settings.


**Instance Name \***  
openstackclient\_jumpstation

**Description**  
jumpstation where openstackclient will be installed

**Availability Zone**  
nova

**Count \***  
1

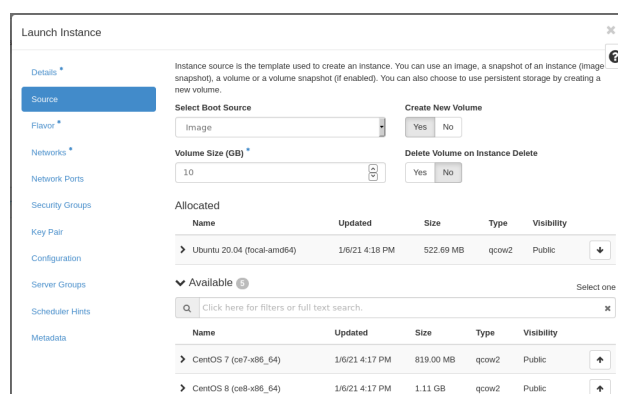
Total Instances (10 Max)  
10%  
0 Current Usage  
1 Added  
9 Remaining

< Back Next > 

The minimal requirements needed to launch an instance are marked by an asterisk.

After the **Details** section is filled out, move on to the **Source** tab on the left. Here you will specify a source to boot from, which is typically going to be an image.

### Choose Source



Launch Instance

Instance source is the template used to create an instance. You can use an image, a snapshot of an instance (image snapshot), a volume or a volume snapshot (if enabled). You can also choose to use persistent storage by creating a new volume.

**Select Boot Source**  
Image

**Volume Size (GB) \***  
10

**Create New Volume**  
Yes No

**Delete Volume on Instance Delete**  
Yes No

**Allocated**

Name	Updated	Size	Type	Visibility
Ubuntu 20.04 (focal-amd64)	1/6/21 4:18 PM	522.69 MB	qcow2	Public

**Available**

Click here for filters or full text search.

Name	Updated	Size	Type	Visibility
CentOS 7 (xe7-x86_64)	1/6/21 4:17 PM	819.00 MB	qcow2	Public
CentOS 8 (ce8-x86_64)	1/6/21 4:17 PM	1.11 GB	qcow2	Public

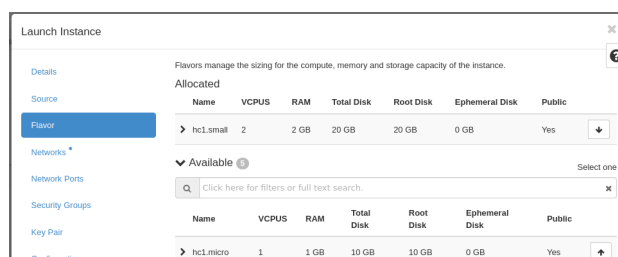
You need to choose a **Boot Source**, specify whether or not to create a volume, and finally the image source.

Here, the **Boot Source** is selected as **Image** and a new Volume of size 10GB will be created. The images are listed on this same form and you can use the up arrow on the right to choose the image.

Next you need to specify a flavor. The flavor is a way to define resource allocation to an instance. The number of vCPUs, RAM, and disk space are defined using flavors.

This example has selected the **hc1.small** flavor.

### Choose Flavor



Launch Instance

Flavors manage the sizing for the compute, memory and storage capacity of the instance.

**Allocated**

Name	VCPUS	RAM	Total Disk	Root Disk	Ephemeral Disk	Public
hc1.small	2	2 GB	20 GB	20 GB	0 GB	Yes

**Available**

Click here for filters or full text search.

Name	VCPUS	RAM	Total Disk	Root Disk	Ephemeral Disk	Public
hc1.micro	1	1 GB	10 GB	10 GB	0 GB	Yes

Next, for the **Network** select the private network created in the previous guide or choose your own private network.

### Choose Network

Launch Instance

Networks provide the communication channels for instances in the cloud.

Select networks from those listed below.

Network	Subnets Associated	Shared	Admin State	Status
Internal	Internal_subnet_1	No	Up	Active

Select at least one network

Click here for filters or full text search.

Network	Subnets Associated	Shared	Admin State	Status
External	Internet	Yes	Up	Active

Here, the **Internal** Network is the network this instance will be associated with. Neutron will obtain an IP from that network and assign it to this instance.

After the **Network** section, you'll want to be sure you have a **Security Group** that matches your needs. This example will use the **default** security group and in addition will use one made previously, called **SSH**.

### Choose Security Group

Launch Instance

Select the security groups to launch the Instance in.

Select one or more

Name	Description
SSH	General SSH security group
default	Default security group

Click here for filters or full text search.

No available items

By default, only the **default** security group will be selected. To set the **SSH** security group, choose the up arrow on the right of the **SSH** security group listing.

Finally, an SSH key pair will be added to the instance.

### Choose SSH Key

Launch Instance

A key pair allows you to SSH into your newly created instance. You may select an existing key pair, import a key pair, or generate a new key pair.

+ Create Key Pair   - Import Key Pair

Allocated

Displaying 1 item

Name	Type
ssh_key	ssh

Displaying 0 items

Select one

Click here for filters or full text search.

Displaying 0 items

No items to display.

Since **ssh\_key** is the only SSH public key associated with this cloud it is the default option and selected for you already. If there were other keys you could add them from the list below. You can create a key pair or upload a public key using this form.

With these steps done, you are ready to create the instance. Click **Launch Instance** to do so.


When the instance has finished being setup and active it will appear this way:

### Instance listing:

# Instances

Instance ID =

Filter

 Launch Instance

Delete Instances

More Actions

Displaying 1 item

<input type="checkbox"/>	Instance Name	Image Name	IP Address	Flavor	Key Pair	Status	Availability Zone	Task	Power State	Age	Actions	
<input type="checkbox"/>	openstackclient_jumpstation	-	192.168.0.228	hcl1_small	st_timespace	Active	us-east-1	nova	None	Running	0 minutes	Create Snapshot

Displaying 1 item

### ***Troubleshooting Instance Error Status***

Sometimes, instance creation will not succeed which can be for a variety of reasons. Running `openstack server list` will show the **Status** of an instance.

Here's an example of an instance with Status, "ERROR":

```
$ openstack server list
```

ID	Name	Status	Networks
35d624fe-785d-4915-aa7e-4bb580b29325	centos_instance_2	ERROR	

Get more information on the error status by running:

```
$ openstack server show $INSTANCE_UUID --fit-width
```

Look for the **fault** row for the reason as to why the instance failed to create.

It is suggested to look over <https://docs.openstack.org/nova/train/admin/support-compute.html> for common issues that may arise and how to troubleshoot them.

### ***Location of OpenStack service logs***

You may need to look at service's logs to get a better idea of an issue. The default log location for an OpenStack deployment is `/var/log/$service_name`, however OpenStack in this case has been deployed using `kolla-ansible` and logs are stored in `/var/log/kolla/$service_name` on each hardware node.

---

### ***Next Steps***

The next guide in this series will explain how to SSH into this instance.

If you wish to skip the SSH guide, see the managing images in OpenStack guide.

### ***Create an OpenStack Instance with OpenStackClient***

The purpose of this guide is to explain how to create an instance using the command line through OpenStackClient.

An example will be demonstrated where an instance is created on a private network. The intent of the instance will be to act as a "jumpstation" where OpenStackClient will be installed. Finally, this jumpstation can be used to administer the OpenStack cloud using OpenStackClient.

An instance is another name for a virtual machine in OpenStack. Instances are created by the Nova service and contribute to the processing power of the cloud.

You'll also learn how to upload or create an SSH key pair, assign storage using a volume, and create a security group. Each of these components will be added to the instance.

**NOTE!** – You will need to have OpenStackClient installed already to follow this guide.

---

## OpenStackClient

OpenStackClient is how you would manage an OpenStack cloud using the command line.

For general information on OpenStackClient and how to install it, see the Day 1 guide.

Throughout these guides OpenStackClient will also be known as OSC.

To get a full list of the available commands, run:

```
openstack --help
```

**NOTE!** – In addition to OpenStackClient, there are other command line clients for various OpenStack services that can be used. For example, [Nova](#) and [Cinder](#) each have their own command line client, as well as other services.

In later releases of OpenStack use of service-specific command line interfaces will be deprecated. When using the command line to administer OpenStack it is recommended to use OpenStackClient where possible as opposed to individual service's command line interfaces such as `nova's CLI`.

---

## Prerequisites for Creating an Instance

Before making an instance, ensure these items exist:

- A flavor
- An image
- A network
- A security group
- An SSH public key

At minimum, a flavor, an image, and a network should exist. You may also want to attach a volume, specify security groups to use, and attach an SSH public key.

---

## SSH key pairs

An SSH key pair will be required to access any instances over SSH. Password authentication is by default disabled in the operating system images.

You can either have an SSH key created or you can upload the public key of your SSH key pair.

### Create an SSH key pair

SSH keys can be managed, created, and uploaded through the command line using OpenStackClient.

To make a key pair, use:

```
$ openstack keypair create KEY_NAME
```

**KEY\_NAME** is the name of the SSH key pair.

This will generate a key pair for you and return the private key. The private key should be kept somewhere safe and be inaccessible to others.

To upload your public key, use:

```
$ openstack keypair create --public-key PATH_TO_PUBLIC_KEY KEY_NAME
```

**KEY\_NAME** is the name of the SSH key pair and **PATH\_TO\_PUBLIC\_KEY** is the path on the filesystem to the public key.

Here's an example of uploading an SSH public key:

```
$ openstack keypair create --public-key ~/.ssh/ssh_key.pub ssh_key_2
```

Field	Value
fingerprint	ff:a4:81:c7:59:07:aa:54:43:39:52:cd:b2:12:aa:fb
name	ssh_key_2
user_id	43317575cccc440fbc38a1f23b45125

## Security groups

A security group in OpenStack controls network access. If no security group is assigned, a default one will be used. The default security group disables all inbound traffic and only allows outbound traffic, which may not be useful.

A security group should be made that allows the type of incoming and outgoing traffic required. OpenStack's default security group does not allow incoming traffic so if you wanted to SSH into an instance, port 22 will be closed. An example of a security group would be one that allows web traffic by opening ports 80 and 443 and you may want another security group that opens port 22 for SSH traffic.

### Create and manage security groups

This section will demonstrate creating a security group that allows SSH traffic from a specific host.

#### Basics

> The command to create a security group appears like so, where **SECURITY\_GROUP** is the name of the security group:

```
$ openstack security group create SECURITY_GROUP
```

> To list security groups, run:

```
$ openstack security group list
```

ID	Name	Description
8639e3c5-47ce-4072-a1f5-1c1e931a8f75	default	Default security group
ebffcf78-52d9-436c-81db-5ea788a0c33d	devstack	
ec8a02ba-4bc2-4b78-a555-902caead87fe	basic_webserver_group	This will open standard ports

> To list the details of a specific security group, where **UUID** is the security group's UUID, run:

```
$ openstack security group show UUID
```

#### Procedure

The following example demonstrates opening port 22 for incoming traffic coming from **50.50.50.50**.

There are two steps to the example: The first step shows creating a security group called **ssh\_demo** and the next step shows adding the SSH rule to that group.

**Open port 22 for incoming SSH traffic:**

```
# Create the security group
$ openstack security group create ssh_demo

# Add rule for port 22
```

```
$ openstack security group rule create --remote-ip 50.50.50.50/32 \
--dst-port 22:22 --ingress --protocol tcp ssh_demo
```

The next section is not required, but shows how to make a security group allowing incoming traffic to the common HTTP ports, 80 and 443.

**Example security group for HTTP incoming traffic:**

```
$ openstack security group rule create --remote-ip 0.0.0.0/0 \
--dst-port 80:80 --ingress --protocol tcp SECURITY_GROUP

$ openstack security group rule create --remote-ip 0.0.0.0/0 \
--dst-port 443:443 --ingress --protocol tcp SECURITY_GROUP
```

## Assign Storage

Storage can be assigned through volumes. A volume in OpenStack is like a removable USB drive that can be attached to instances as seen fit. Cinder is the OpenStack service that allows volumes to be created. A volume can also be used to boot an instance.

### Create and attach a volume

Use the following command to create a volume with size 15GB:

```
$ openstack volume create volume_1 --size 15
```

List the newly created volume:

```
$ openstack volume list
```

ID	Name	Status	Size	Attached to
1fe4f750-62fc-4c53-b661-44481d0f72b3	volume_1	available	15	

Attach the volume to an instance:

```
$ openstack server add volume $INSTANCE_UUID $VOLUME_UUID
```

## Create an instance

If at this step and you have followed the previous steps, you should have everything needed to create the instance. Again, this example will demonstrate how to create an instance on a private network.

**NOTE!** – If your infrastructure does not need to be on a public network, then ensure the appropriate parts are created on a private network, and be sure the security groups are set to only allow the needed traffic through. This information goes a long way in improving overall security when using OpenStack.

### Before creating an instance

Before creating an instance, you will need to collect some information.

**Needed details:**

- The flavor, image, and network UUIDs
- SSH key
- Security group

You can list the above using OpenStackClient. The next section will provide examples using each command and their outputs.

**NOTE!** – When listing items using OpenStackClient, almost everything will have a UUID. However, if a UUID does not exist, you'll need to use the name specified in the `ID` column or the name of the item itself.

Here are commands that can be used to collect the needed information:

#### Summary of commands:

```
$ openstack flavor list
$ openstack image list
$ openstack network list
$ openstack security group list
$ openstack keypair list
```

#### Detailed commands:

##### List flavors:

```
$ openstack flavor list
```

ID	Name	RAM	Disk	Ephemeral	VCPUs	Is Public
hcl.large	hcl.large	8192	80	0	8	True
hcl.medium	hcl.medium	4096	40	0	4	True
hcl.micro	hcl.micro	1024	10	0	1	True
hcl.small	hcl.small	2048	20	0	2	True
hcl.xlarge	hcl.xlarge	16768	160	0	16	True
hraml.medium	hraml.medium	16768	40	0	4	True

Flavors are a way to set things like the amount of RAM, disk space, and CPU cores to be assigned to an instance. The `Is Public` heading says that this flavor is shared among all OpenStack projects.

##### List images:

```
$ openstack image list
```

ID	Name	Status
22c437b3-18cd-4af0-bd3d-ad26c85bb00f	CentOS 7 (ce7-x86_64)	active
cd18f302-d3bf-49a1-8f5d-a6805404b9ff	CentOS 8 (ce8-x86_64)	active
507ca3c5-ed57-4997-b21d-bc0f32406322	Debian 10 (buster-amd64)	active
c19c4d9d-59c9-49e4-8e5f-988d3c1ae4d4	Debian 9 (stretch-amd65)	active
46b14073-bd4e-4aa2-b901-63c86ab15961	Ubuntu 18.04 (bionic-amd64)	active
40b09133-a1a0-4eal-a8aa-958541314ac5	Ubuntu 20.04 (focal-amd64)	active

Images are typically an operating system image and usually can be found publicly available. CentOS has a repository of [cloud instances](#) for example.

Images can also be created using a snapshot of an instance.

##### List networks:

```
$ openstack network list
```

ID	Name
----	------

0db14de2-9dd0-469f-a76c-46b99119d607	External
bdb2dbfc-e3ef-4da0-a79b-fe48c4a413fb	Internal
fcd8166c-0fdb-4bc9-ab29-ce54a8bd5959	HA network tenant b9e8639372014c0b85cbfaffa6e1b5a8

**List security groups:**

```
$ openstack security group list
```

ID	Name	Description	Project
06785fc8-492c-477b-9928-9708cb550a14	ssh_demo	ssh_demo	b9e8639372
9e20fbef-9eca-4029-8111-625945409ae2	default	Default security group	6a654535b8
cc25e2d6-6a5c-469f-ba89-c2e04e6f6850	SSH	General SSH security group	b9e8639372

**List SSH key pairs:**

```
$ openstack keypair list
```

Name	Fingerprint
st_timespace	a7:ab:bd:9c:78:85:e1:a1:c4:07:0f:6d:e9:36:0b:68

**Procedure to create an instance**

Base command needed to make an instance:

```
$ openstack server create
```

From this command you will specify additional flags such as the SSH key, the network UUID, the instance name, and the like.

For the full list of options to make an instance, run:

```
$ openstack help server create
```

The command to create an instance will look something like this:

```
$ openstack server create --image IMAGE_UUID \
--flavor hcl.medium --network NETWORK_UUID \
--key-name SSH_KEY_NAME --security-group SECURITY_GROUP_UUID \
INSTANCE_NAME
```

Note that IMAGE\_UUID is the UUID of the image you want to use, NETWORK\_UUID is the UUID of the network to be associated with the instance, SSH\_KEY\_NAME is the name of the SSH key, SECURITY\_GROUP\_UUID is the UUID of the security group to use, and INSTANCE\_NAME is the name to call the instance.

**NOTE** – By default, the instance creation will occur in the background. You can add `--wait` to the flags to have the command wait until the instance creation is done which will show you the status of instance creation.

**Example instance creation:**

Here are the collected details to make an instance from the previous section:

- Flavor: hc1.small
- Image UUID: 40b09133-a1a0-4ea1-a8aa-958541314ac5
- Network UUID: bdb2dbfc-e3ef-4da0-a79b-fe48c4a413fb
- SSH key: st\_timespace
- Security group UUID: cc25e2d6-6a5c-469f-ba89-c2e04e6f6850

**Full command:**



```
$ openstack server create --image 40b09133-ala0-4ea1-a8aa-958541314ac5 \
--flavor hcl.small --network bdb2dbfc-e3ef-4da0-a79b-fe48c4a413fb \
--key-name st_timespace --security-group \
cc25e2d6-6a5c-469f-ba89-c2e04e6f6850 openstackclient_js_demo
```

After creating the instance, verify the build process by running `$ openstack server show INSTANCE_NAME`.

The `status` column will indicate the status. See the [Nova compute API documentation](#) for a list of instance status meanings and additional commands that can be used to troubleshoot any issues.

You also may find the table output is too large for your display. Add `--fit-width` to the command to force the table to fit to your display's width making reading easier.

## Troubleshooting

**NOTE!** – This section is incomplete but will receive updates in the future.

There may be an issue with spawning the instance. Using the above command will display any errors with the build process under the `fault` field.

## Troubleshooting Instance Error Status

Sometimes, instance creation will not succeed which can be for a variety of reasons. Running `openstack server list` will show the **Status** of an instance.

Here's an example of an instance with Status, "ERROR":

```
$ openstack server list
```

ID	Name	Status	Networks
35d624fe-785d-4915-aa7e-4bb580b29325	centos_instance_2	ERROR	

Get more information on the error status by running:

```
$ openstack server show $INSTANCE_UUID --fit-width
```

Look for the **fault** row for the reason as to why the instance failed to create.

It is suggested to look over <https://docs.openstack.org/nova/train/admin/support-compute.html> for common issues that may arise and how to troubleshoot them.

## Location of OpenStack service logs

You may need to look at service's logs to get a better idea of an issue. The default log location for an OpenStack deployment is `/var/log/$service_name`, however OpenStack in this case has been deployed using `kolla-ansible` and logs are stored in `/var/log/kolla/$service_name` on each hardware node.

The next guide in this series will explain how to connect an instance to the provider network, allowing Internet access.

## Next Steps

The next guide in this series will explain how to SSH into this instance.

If you wish to skip the SSH guide, see the managing images in OpenStack guide.

## SSH into an Instance

This guide serves to explain how you can SSH into an instance.

The OpenStackClient can be used to SSH into instances. Instances can be created on both public and private networks. The location from where the SSH command is being issued has to be able to access the instance over the network. Typically instances created either on the provider network or that have been assigned a floating IP can be accessed from any machine that is connected to the Internet.

### Access Instance associated with a Public IP

If the instance has a floating IP or is on the provider network, then the instance can be accessed by any machine that has OpenSSH installed or from any machine that has OpenStackClient installed.

#### Using SSH

To SSH into an instance, the machine you connect from has to be able to connect to the SSH port (typically 22) of the instance. This means the security group associated with the instance must allow SSH traffic from the machine you intend to SSH from. See the create an instance guide for how to create security groups.

An example command to SSH into an instance:

```
$ ssh -i ~/.ssh/KEY USER@50.50.50.50
```

In this example, `~/.ssh/KEY` is your private key, **USER** is the SSH username, and **50.50.50.50** is an IP you can connect to from the machine you intend to SSH from.

#### Using OpenStackClient

It is also possible to use OpenStackClient to SSH into a machine. To be able to use OpenStackClient for this purpose, the machine you connect from must have SSH access to the instance. Again, ensure the instance has a security group that allows SSH access from the host you intend to connect with.

Base command to SSH into an instance:

```
$ openstack server ssh
```

Example command to SSH into an instance which has been created on the provider network:

```
$ openstack server ssh --login centos --identity ~/.ssh/ssh_key --address-type fixed cf491b0
```

`--address-type` can be public, private, or fixed

A **fixed** `--address-type` means the IP assigned to the instance is a static IP. When an instance is created on the public network, a fixed IP will be assigned to it.

### Access Instance associated with a Private IP

It is also possible to SSH into an instance that is on a private network. This will have to be done from one of the hardware nodes which has to be associated with the appropriate private network. The private key of the SSH key pair should be on that node as well.

Listed is the instance in question to connect to:

```
$ openstack server list
```

ID	Name	Status	Networks

```
| e93b3344-6d78-4273-880f-220b7fbec417 | test_5 | ACTIVE | Internal=19
+-----+-----+-----+-----+
```

It can be seen the IP associated with it is on a private network.

Determine what compute node the instance is on:

```
$ openstack server show e93b3344-6d78-4273-880f-220b7fbec417
+-----+-----+
| Field | Value |
+-----+-----+
| OS-DCF:diskConfig | AUTO |
| OS-EXT-AZ:availability_zone | nova |
| OS-EXT-SRV-ATTR:host | hcl.example.com |
| OS-EXT-SRV-ATTR:hypervisor_hostname | hcl.example.com |
+-----+-----+
```

Note that some of the output is truncated.

Next, SSH into the appropriate compute node and then find the appropriate private network. This can be done by listing the network namespaces on that node.

List network namespaces:

```
# ip netns
qrouter-4dc1debc-ecf3-42e1-89c6-e2b99fc2c3dd (id: 0)
qdhcp-a54fc8a3-89b1-4ec3-a441-79c6cfe0e915 (id: 3)
qdhcp-55d31bd5-77ba-4ed0-ab6e-99554b33aa90 (id: 1)
```

List the interfaces for a network namespace using the format `ip netns exec $network_namespace $command`, so for example:

```
# ip netns exec qdhcp-a54fc8a3-89b1-4ec3-a441-79c6cfe0e915 ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
164: tapaa57977f-ca: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UNKNOWN
    link/ether fa:16:3e:a3:db:34 brd ff:ff:ff:ff:ff:ff
    inet 192.168.0.2/24 brd 192.168.0.255 scope global tapaa57977f-ca
        valid_lft forever preferred_lft forever
    inet 169.254.169.254/16 brd 169.254.255.255 scope global tapaa57977f-ca
        valid_lft forever preferred_lft forever
    inet6 fe80::f816:3eff:fea3:db34/64 scope link
        valid_lft forever preferred_lft forever
```

From here the output shows the subnet 192.168.0.2/24 listed and 192.168.0.186 is on that network, so to SSH into this instance, an example command would take a form like this:

```
# ip netns exec qdhcp-a54fc8a3-89b1-4ec3-a441-79c6cfe0e915 ssh -i $ssh_key centos@192.168.0.2
Activate the web console with: systemctl enable --now cockpit.socket

Last login: Mon Aug 24 21:59:47 2020 from 192.168.0.2
[centos@test-5 ~]$
```

## Next Steps

The next guide explains how to manage images in OpenStack.

## OpenStack Images

Images in OpenStack are what powers instances. This guide will explain how to use images that are already uploaded into [Glance](#), the OpenStack service responsible for images and how to upload your own.

An image is a file that contains a bootable operating system. Many different cloud image sources are available for download from major operating system providers like CentOS, Ubuntu, and Debian to name a few. You can also make your own images from [scratch](#) or create them from volumes or running instances. Snapshots of instances can be created which can serve both as a backup and also a template for other instances.

It is possible to create your own images however this process is outside the scope of this guide. For more information on how to do this, see the [documentation](#) from OpenStack.

### List and upload images

This section will cover how to view and upload images into OpenStack. Either Horizon or the command line using OpenStackClient can be used.

#### Using Horizon

In Horizon, the images can be found on the left within the **Compute** tab under **Images**. From this interface, you could launch an instance using one of the images and also upload your own.

#### List images

Images in Horizon:

Owner	Name *	Type	Status	Visibility	Protected	Disk Format	Size	
admin	CentOS 7 (x86_64)	Image	Active	Public	No	QCOW2	819.00 MB	Launch
admin	CentOS 8 (x86_64)	Image	Active	Public	No	QCOW2	683.00 MB	Launch

### Upload and create images

In the same Horizon interface, you can upload an image, make your own or make an image snapshot from a pre-existing instance or volume.

Many operating system providers have cloud images available for download already. For example, see the [image repository](#) from CentOS.

Similarly, other major Linux operating system creators provide images that can be used in OpenStack.

To upload an image, login to Horizon, navigate to **Compute**, then **Images**. From there, click the **Create Image** button. A form will appear where you can fill out the needed details. Click the **?** icon for assistance with specifics.

Upload an image:

This example will walk through uploading the [CirrOS](#) image.

Ensure the image to upload has been downloaded to your machine. CirrOS has been downloaded from the [CirrOS download page](#). The [latest version](#) is 0.5.1. Be sure to download the latest version of the image needed.

With the image on your machine, the form details can be filled.

- **Image Details**

- **Image Name** - specify the name for the uploaded image
- **Image Description** - option field used to describe the image

- **Image Source**

- **File** - this is the location of the image file on your machine
- **Format** - the format should be **QCOW2 - QEMU Emulator**

Optionally, you can specify requirements, like minimum disk space and RAM, for the image under the **Image Requirements** heading.

## Using the command line

You can also view and upload images using OpenStackClient over the command line.

### List images

To view available images, use:

```
$ openstack image list
```

ID	Name	Status
6986aaf9-e602-4307-9c4f-6377795ff890	CentOS 7 (ce7-x86_64)	active
d3f027f2-730f-4794-9288-44ed5c69050d	CentOS 8 (ce8-x86_64)	active
29c5d143-9fda-4f8e-b364-c2dc29851101	Debian 10	active
6eabc356-c373-4260-8f6f-7d83ff80df21	Snap-1	active
994b6e02-b8c4-4652-80bf-3d56c92b4b19	Ubuntu 18.04 (bionic-amd64)	active
20b56b2d-1af5-46e1-a5c6-fa1f9a45245d	Ubuntu 19.10 (eoan-amd64)	active
d35559ff-c9fd-4c18-be8b-1a74ecaleb38	Ubuntu 20.04 (focal-amd64)	active

To get more information about an image, run:

```
$ openstack image show ID
```

where **ID** can be the Name or the ID column in the above output.

Example showing the details of an image:

```
$ openstack image show fa8eb9bd-9ccc-4d3f-b87b-6edb5450a57a --fit-width
```

Field	Value
checksum	1d3062cd89af34e419f7100277f38b2b
container_format	bare
created_at	2020-09-09T20:50:25Z
disk_format	qcow2
file	/v2/images/fa8eb9bd-9ccc-4d3f-b87b-6edb5450a57a/file
id	fa8eb9bd-9ccc-4d3f-b87b-6edb5450a57a
min_disk	0
min_ram	0
name	cirros
owner	5ad1f9e795604f4390d274d7388c4b9f

The output of the above command has been truncated.

### Upload an image

To upload an image, run:

```
openstack image create
```

For assistance with the CLI options, run:

```
openstack help image create
```

OpenStack provides an operating system called [CirrOS](#) that has the minimum requirements to be an operating system that is generally used to test instance creation.

To upload this CirrOS image into the glance service, you'll need to first download it to where the OpenStackClient lives, then the base command `openstack image create` can be used to upload the image.

#### Steps to upload the CirrOS image:

Obtain the latest CirrOS image:

```
$ wget http://download.cirros-cloud.net/0.5.1/cirros-0.5.1-x86_64-disk.img
```

Upload the image into glance:

```
$ openstack image create cirros --container-format bare --disk-format \
qcow2 --file PATH_TO_CIRROS_IMAGE
```

Note that **PATH\_TO\_CIRROS\_IMAGE** should be path to the image file.

List the newly uploaded image:

```
$ openstack image list
```

ID	Name	Status
fa8eb9bd-9ccc-4d3f-b87b-6edb5450a57a	cirros	active

## Create Image Snapshot

### Snapshot using Horizon

An image can also be created from an existing instance by making use of the image snapshot feature.

In the dashboard, look for **Compute**, then **Instances**. From there, select an instance, and click **Create Snapshot** from the drop down on the right.

Create image snapshot:

Instance Name	Image Name	IP Address	Flavor	Key Pair	Status	Availability Zone	Task	Power State	Age	Actions
image volume testing	CentOS 8 (je 8-x86_64)	10.0.0.121	hcl_small	ssh_key	Active	nova	None	Running	23 hours, 38 minutes	Create Snapshot
debian_test	Debian 10	192.168.0.178	hcl_small	ssh_key	Active	nova	None	Running	2 days, 1 hour	Create Snapshot

This will create an image of that instance and the state of the disk will be preserved. You can use this to launch other instances. Should you want to launch an instance using this snapshot, it will be listed in the **Launch Instance** form under the **Select Boot Source** dropdown as a volume snapshot.

Create snapshot form:

Create Instance form showing the volume snapshot recently created:

### Snapshot Using the command line

A snapshot of an instance can be made using the command line with OpenStackClient.

The command to create a snapshot is:

```
$ openstack server image create --name SNAPSHOT_NAME INSTANCE_NAME
```

- **SNAPSHOT\_NAME** is the name to call the snapshot
- **INSTANCE\_NAME** is the name of the instance the snapshot is created from. The instance UUID could also be used.

After creating a snapshot verify success by:

- `openstack image list`
- `openstack image show SNAPSHOT_UUID` - SNAPSHOT\_UUID refers to the UUID of the snapshot made.

### Create and verify an instance snapshot

Make a snapshot of the **debian\_test** instance:

```
$ openstack server image create --name debian_snapshot debian_test --fit-width
```

Field	Value
checksum	d41d8cd98f00b204e9800998ecf8427e
container_format	bare
created_at	2020-11-05T16:21:25Z
disk_format	qcow2
file	/v2/images/161eb918-f470-4bb8-9572-43ee9eb27ceb/file
id	161eb918-f470-4bb8-9572-43ee9eb27ceb
min_disk	20
min_ram	0
name	debian_snapshot
owner	5ad1f9e795604f4390d274d7388c4b9f

List available snapshots:

```
$ openstack image list --fit-width
```

ID	Name	Status
6986aaf9-e602-4307-9c4f-6377795ff890	CentOS 7 (ce7-x86_64)	active
d3f027f2-730f-4794-9288-44ed5c69050d	CentOS 8 (ce8-x86_64)	active
29c5d143-9fda-4f8e-b364-c2dc29851101	Debian 10	active
161eb918-f470-4bb8-9572-43ee9eb27ceb	debian_snapshot	active

The snapshot created is listed as the last entry under ID of **161eb918-f470-4bb8-9572-43ee9eb27ceb**.

List details of the snapshot:

```
$ openstack image show 161eb918-f470-4bb8-9572-43ee9eb27ceb --fit-width
```

Field	Value
checksum	d41d8cd98f00b204e9800998ecf8427e
container_format	bare
created_at	2020-11-05T16:21:25Z
disk_format	qcow2
file	/v2/images/161eb918-f470-4bb8-9572-43ee9eb27ceb/file
id	161eb918-f470-4bb8-9572-43ee9eb27ceb
min_disk	20
min_ram	0
name	debian_snapshot
owner	5ad1f9e795604f4390d274d7388c4b9f

## Ceph, Cinder and Swift

Ceph, Cinder, and Swift are used in OpenStack for data storage. Ceph is the backend storage software used and is not an OpenStack service, but rather the backend system where data for Cinder and Swift is stored. Cinder is OpenStack's service for block storage and Swift is the service responsible for object storage.

The purpose of this guide is to briefly explain each of these services and how they relate to one another.

### Ceph

Ceph is a unified, distributed storage system designed for excellent performance, reliability and scalability and is the backend storage used by this deployment of OpenStack.

The Octopus [release](#) of Ceph is currently being used.



## ***Ceph Data Durability***

Ceph has different strategies for data durability, or data redundancy. See Ceph's [documentation](#) for information on how Ceph stores data.

Typically a Ceph configuration will have three replicas of data per pool as this is the best tradeoff between performance and data durability. If data durability and cost are not as important, Ceph can be configured with two OSDs per pool, with the tradeoff being less physical hardware used to store data, meaning less cost. Another option in Ceph for replicating data is called [erasure coding](#). This method is a software approach to replicating data that requires less disk space than using three replica OSDs at the expense of greater CPU and RAM usage.

---

## ***Cinder***

Cinder is a block storage service for OpenStack which allows volumes to be created. A volume is a detachable storage device and can only be attached to one instance at a time.

This service virtualizes the management of block storage devices. End users make use of the storage service without needing to know where the storage is deployed or on what type of device it is stored. Drivers can be created which allow for multiple storage backend solutions, including local via LVM (Logical Volume Manager), Ceph, and others.

## ***More Information***

For more information about Cinder, check out these resources.

- [Cinder OpenStack wiki](#)
- [Cinder's OpenStack documentation](#)

## ***What is block storage?***

Block storage is strategy of storing data where the data is stored in fixed size chunks called blocks. The address of where the block is located is the only metadata associated with each block.

---

## ***Swift***

Swift is the OpenStack service that handles object storage. This service allows containers to be created that can store objects, for example a collection of images, audio files, or a combination of any sort of files. Files can be uploaded through the Swift API or using Horizon.

Swift uses Ceph's [RadosGW](#) as a means to handle object storage, effectively meaning Ceph is used to store Swift data.

## ***More Information***

For more information about Swift, check out these resources.

- [Swift OpenStack wiki](#)
- [Swift's OpenStack documentation](#)

## ***What is object storage?***

Object storage is a strategy of storing data in various formats of varying sizes. Objects are stored as entire objects, so an mp3 file is stored completely as that along with any metadata associated with that object. The problem object storage solves is data scalability.

### *How do Cinder and Swift relate to Ceph?*

Both Cinder and Swift can use different storage backends, ranging from local disks or a storage cluster like Ceph.

This deployment of Cinder and Swift both use Ceph as a backend for data storage.

Swift has been configured to use Ceph's Object Gateway also known as **RADOS Gateway**.

### *Using the Ceph CLI for Common Operational Tasks*

This guide will go over some of the common tasks to do using the Ceph's CLI tool.

---

#### *Common tasks*

Get status of ceph cluster:

```
$ ceph status
```

List ceph pools:

```
$ ceph osd lspools
```

List overall disk usage:

```
$ ceph df
```

Get status of all OSDs:

```
$ ceph osd status
```

Get status of monitors:

```
$ ceph mon stat
```

Get status of placement groups (PGs):

```
$ ceph pg stat
```

List amount of disk usage per OSD:

```
$ ceph osd df
```

Change number of replicated OSDs for a pool:

```
$ ceph osd pool set $POOL size $COUNT
```

#### *More detailed ``ceph status`` output*

Get Ceph cluster status using `ceph status`:

```
# ceph status
cluster:
  id:      228c6e47-b60e-4710-9bcb-54316f752857
  health: HEALTH_WARN
          1 pools have many more objects per pg than average
          clock skew detected on mon.hc2, mon.hc3

services:
  mon: 3 daemons, quorum hc1,hc2,hc3 (age 21h)
  mgr: hc2(active, since 8d), standbys: hc1, hc3
```

```

osd: 3 osds: 3 up (since 21h), 3 in (since 22h)
rgw: 3 daemons active (hc1.rgw0, hc2.rgw0, hc3.rgw0)

task status:

data:
  pools:    13 pools, 225 pgs
  objects:  15.65k objects, 59 GiB
  usage:    169 GiB used, 2.5 TiB / 2.6 TiB avail
  pgs:      225 active+clean

io:
  client:   29 KiB/s rd, 34 op/s rd, 0 op/s wr

```

## Ceph Replication, Compression, and Erasure Coding

Ceph is naturally resilient to data loss. It accomplishes this by replicating the same set of data several times. In other words, Ceph stores multiple copies of the same set of data. Data is either replicated across several Object Storage Daemons (OSDs) or is replicated using [erasure coding](#).

This guide will explain a bit how Ceph has been configured for Private Cloud Core and also cover adjusting the number of replicated OSDs per pool.

---

### Ceph Replication

With Private Cloud Core, Ceph has been distributed across each hard drive on each hardware node and the data in each pool is being replicated across three OSDs.

The configuration is put into place using `ceph-ansible` and the repository for this software is located in [Github](#). Use of `ceph-ansible` assumes a base understanding of how Ansible works.

Ceph has a concept of pools and there are typically several pools that store different sets of data. Each pool can be configured with different replication options, such as erasure coding for one or some pools and other pools can be set to have multiple OSDs. More on the options for configuring pools can be found in Ceph's [documentation](#).

Another way to replicate data in Ceph is to use erasure coding. This is a software approach to data replication that uses less disk space, but has more overhead in terms of CPU and RAM usage. See the InMotion Hosting [support center](#) for more information on this subject.

With erasure coding, this replication type must be set upon creation of the Ceph pool and cannot be changed after a pool has been created. If you have a set of data replicated using erasure coding and want to change how data is replicated, a pool with the replication strategy intended needs to be created then the data would be migrated to that pool.

Here's an example of the pools you may see in the Ceph cluster using the `ceph osd lspools ceph` CLI command:

```

$ ceph osd lspools
1 device_health_metrics
2 images
3 volumes
4 vms
5 backups
6 metrics
7 manila_data
8 manila_metadata
9 .rgw.root
10 default.rgw.log
11 default.rgw.control
12 default.rgw.meta
13 default.rgw.buckets.index

```

See the [documentation from Ceph](#) for details on how Ceph handles data durability, otherwise known as data replication.

---

### **Setup replication across 2 OSDs**

Let's say you wanted the `volumes` pool to have replication between two OSDs. The command to set that up would look like this:

```
$ ceph osd pool set volumes size 2
set pool 3 size to 2
```

### **Setup replication across 3 OSDs**

The recommended number of replicated OSDs is three and to set that for the `volumes` pool, this command needs to be used:

```
$ ceph osd pool set volumes size 3
set pool 3 size to 3
```

### **Erasure Coding**

For information on erasure coding in general, see the InMotion Hosting [support center](#).

## **OpenStack CLI for Common Operational Tasks**

### **Background**

OpenStackClient (OSC) is the name of a command line interface for OpenStack which can be used to administer an OpenStack cloud. The same functionality found in the Horizon interface can also be found using OpenStackClient.

---

### **Get started**

OpenStackClient will need to be installed before it can be used. See this guide for installation instructions.

For a full list and explanation of the available options use `openstack help` or refer to the [OpenStack documentation](#).

Additional command line applications exist for some services such as Nova and Cinder but will eventually be deprecated. While these command line utilities can be used, it is recommended that the OpenStack cloud be administered using solely the OSC. The goal of OpenStack client is to provide all the needed commands to administer an OpenStack cloud under one application.

---

### **Output formatting**

Typically when OpenStackClient commands are issued, a table of formatted data is returned. When doing batch operations you may want to extract just the **UUID** of an item. It is possible to have OpenStackClient return exactly the information you need which can be useful for scripts or running the same action on multiple items.

For example:

```
output formatters:
  output formatter options

  -f {csv,json,table,value,yaml}, --format {csv,json,table,value,yaml}
                                the output format, defaults to table
  -c COLUMN, --column COLUMN
```

```

        specify the column(s) to include, can be repeated to
        show multiple columns
--sort-column SORT_COLUMN
        specify the column(s) to sort the data (columns
        specified first have a priority, non-existing columns
        are ignored), can be repeated

```

Say you want just the **UUID** of all the servers on a host. You can run something like this to achieve that goal (where `example_host` is the host you are working with):

```
$ openstack server list --host example_host -f value -c ID
```

## Common Tasks

Below is a list of common operational tasks that can be done with the OpenStack Client. This guide will include examples of how to perform the following tasks:

- Manage OpenStack users, including listing, creating, updating and removing users
- Manage instances by creating them, stopping and starting them, creating a snapshot, plus much more.
- Live migration of instances
- Troubleshoot instance issues
- Upload images
- Create a network
- Create security groups
- Manage SSH key pairs
- Collect details about OpenStack environment

## Manage OpenStack users

In OpenStack there exists the admin user account which has the ability to create additional users. Typically the admin account is used only when that level of privilege is needed otherwise individual user accounts should be used when interacting with an OpenStack cloud.

The following commands can be used to list, create, update, and remove OpenStack users:

List users:

```
$ openstack user list
```

Create a user where **PROJECT\_NAME** is the name of the project, **PASSWORD** is the password to set, and **USERNAME** is the username:

```
$ openstack user create --project PROJECT_NAME --password PASSWORD \
  USERNAME
```

The base command to update a user is `openstack user set USERNAME` where `USERNAME` is the username in question.

Using that base command, you can enable or disable a user account or change details, such as the email address associated with the user.

Disable a user:

```
$ openstack user set USERNAME --disable
```

Enable a user:

```
$ openstack user set USERNAME --enable
```

Change the email address, where **EMAIL\_ADDRESS** should be the email to set:

```
$ openstack user set USERNAME --email EMAIL_ADDRESS
```

## Instance Management

### Create an instance

See the Day 1 guide for information on how to create an instance.

### Stop and start an instance

Stop an instance:

```
$ openstack server stop
```

Start an instance:

```
$ openstack server start
```

### Create an instance snapshot

Here's an example:

```
$ openstack server image create --name SNAPSHOT_NAME INSTANCE_NAME
```

You can verify the snapshot has been created by using both `openstack image list` to find the newly created snapshot, then `openstack image show SNAP_SHOT_UUID` to get details on that snapshot.

## Perform live migration of instances

Sometimes it is necessary to migrate instances from one compute node to another if for example a compute node needs maintenance. It is possible to live migrate instances before bringing down that node.

Before you begin, you'll need the server UUID, the host that server is currently running on and the host to migrate to.

List servers:

```
$ openstack server list
```

ID	Name	Status	Networks
8de97c6f-bd4e-4f23-b8ee-1d9841082760	test_7	ACTIVE	Internal=19
120fc769-ec99-4025-b456-320d8a17a158	test_6	ACTIVE	Internal=19
e93b3344-6d78-4273-880f-220b7fbec417	test_5	ACTIVE	Internal=19

Get the host of the `test_5` server:

```
$ openstack server show e93b3344-6d78-4273-880f-220b7fbec417
```

Field	Value
OS-DCF:diskConfig	AUTO
OS-EXT-AZ:availability_zone	nova
OS-EXT-SRV-ATTR:host	hcl.example_host

The above output of `openstack server show` is truncated.

Find the compute node to migrate to:

```
$ openstack compute service list
```

ID	Binary	Host	Zone	Status	State	Updated At
18	nova-scheduler	hc1.example_host	internal	enabled	up	2020-09-14T21:...
57	nova-scheduler	hc2.example_host	internal	enabled	up	2020-09-14T21:...
66	nova-scheduler	hc3.example_host	internal	enabled	up	2020-09-14T21:...
3	nova-conductor	hc1.example_host	internal	enabled	up	2020-09-14T21:...
9	nova-conductor	hc2.example_host	internal	enabled	up	2020-09-14T21:...
15	nova-conductor	hc3.example_host	internal	enabled	up	2020-09-14T21:...
24	nova-compute	hc3.example_host	nova	enabled	up	2020-09-14T21:...
27	nova-compute	hc1.example_host	nova	enabled	up	2020-09-14T21:...
30	nova-compute	hc2.example_host	nova	enabled	up	2020-09-14T21:...

You can either choose the host to migrate to or one could be automatically selected. The destination host should have a functioning nova-compute service running.

Before migrating to a specific host, ensure the host has enough resources:

```
$ openstack host show hc2.example_host
```

Host	Project	CPU	Memory MB	Disk GB
hc2.example_host	(total)	8	64243	2682
hc2.example_host	(used_now)	0	10240	0
hc2.example_host	(used_max)	0	0	0

Perform the migration to a specific host:

```
$ openstack --os-compute-api-version 2.79 server migrate \
--live-migration --host hc2.example_host \
e93b3344-6d78-4273-880f-220b7fbec417
```

Note that to use the `--host` flag, the Nova API version must be specified using `--os-compute-api-version`. In this case, the maximum version that can be used is 2.79. More on the Nova API version release history is [here](#).

Running `openstack server list` should show the instance status as `MIGRATING`:

```
$ openstack server list
```

ID	Name	Status	Networks
8de97c6f-bd4e-4f23-b8ee-1d9841082760	test_7	ACTIVE	Internal=
120fc769-ec99-4025-b456-320d8a17a158	test_6	ACTIVE	Internal=
e93b3344-6d78-4273-880f-220b7fbec417	test_5	MIGRATING	Internal=

Confirm the instance has been migrated to the destination host:

```
$ openstack server show e93b3344-6d78-4273-880f-220b7fbec417
```

Field	Value
OS-DCF:diskConfig	AUTO
OS-EXT-AZ:availability_zone	nova
OS-EXT-SRV-ATTR:host	hc2.example_host
OS-EXT-SRV-ATTR:hypervisor_hostname	hc2.example_host
OS-EXT-SRV-ATTR:instance_name	instance-0000008d
OS-EXT-STS:power_state	Running
OS-EXT-STS:task_state	None
OS-EXT-STS:vm_state	active

OS-SRV-USG:launched_at	2020-08-24T21:51:42.000000
OS-SRV-USG:terminated_at	None
accessIPv4	
accessIPv6	
addresses	Internal=192.168.0.186
config_drive	
created	2020-08-24T21:51:16Z
flavor	hcl.small (hcl.small)
hostId	a4a089ceb3b247eeb47dde6f58ed85444cf18bf763453e7fbf77
id	e93b3344-6d78-4273-880f-220b7fbec417
image	CentOS 8 (ce8-x86_64)
key_name	nw_1
name	test_5
progress	0
project_id	5ad1f9e795604f4390d274d7388c4b9f
properties	
security_groups	name='basic_webserver_group'
status	ACTIVE
updated	2020-09-15T21:53:49Z
user_id	43317575cccc440fbc38a1f23b45125
volumes_attached	

## Troubleshoot instance issues

### Note

This section will be filled out as common scenarios occur.

## Upload images

To create an image, use:

```
openstack image create
```

For assistance with the CLI options, use:

```
openstack help image create
```

OpenStack provides an operating system called [Cirros](#) that has the minimum requirements to be an operating system that is generally used to test instance creation.

This example will explain how to upload the Cirros image into glance.

To get this Cirros image into the glance service, you'll need to first download it to where the OpenStackClient lives, then the base command `openstack image create` can be used to upload the image.

Download Cirros and upload into glance:

```
# grab the latest Cirros image
$ wget http://download.cirros-cloud.net/0.5.1/cirros-0.5.1-x86_64-disk.img

# upload the image into glance
$ openstack image create cirros --container-format bare --disk-format
qcow2 --file PATH_TO_CIRROS_IMAGE
```



List the newly uploaded image:

```
$ openstack image list
```

ID	Name	Status
fa8eb9bd-9ccc-4d3f-b87b-6edb5450a57a	cirros	active

## Create a private network

Listed are the steps needed to create a private network and connect it to the provider network. Variables are presented in all capital and should be replaced accordingly.

Create a network:

```
$ openstack network create NETWORK_NAME
```

Create a subnet on that network:

```
$ openstack subnet create --subnet-range 10.0.0.0/24 --network NETWORK_NAME
SUBNET_NAME
```

Create a router that will connect to an external, public-facing network:

```
$ openstack router create ROUTER_NAME
```

Add the subnet to the router:

```
$ openstack router add subnet ROUTER_NAME SUBNET_NAME
```

Add the external network gateway:

```
$ openstack router set --external-gateway EXTERNAL_NETWORK_UUID \
ROUTER_NAME
```

## Create security groups

Here's an example that opens inbound traffic for all IPs on ports 80 and 443.

Create a security group where **SECURITY\_GROUP** is the name of the security group:

```
$ openstack security group create SECURITY_GROUP
```

List security groups:

```
$ openstack security group list
```

ID	Name	Description
8639e3c5-47ce-4072-alf5-1cle931a8f75	default	Default security group
ebffcf78-52d9-436c-81db-5ea788a0c33d	devstack	
ec8a02ba-4bc2-4b78-a555-902caead87fe	basic_webserver_group	This will open standard por

List details of a security group:

```
$ openstack security group show UUID
```

Open ports 80 and 443 on all IP ranges for ingress TCP traffic:

```
$ openstack security group rule create --remote-ip 0.0.0.0/0 \
--dst-port 80:80 --ingress --protocol tcp SECURITY_GROUP

$ openstack security group rule create --remote-ip 0.0.0.0/0 \
--dst-port 443:443 --ingress --protocol tcp SECURITY_GROUP
```

## Add SSH public key

It is recommended an SSH public key be uploaded and this is possible through the OSC.

Running `openstack keypair create KEY_NAME`, where `KEY_NAME` is the name of the SSH key pair, is sufficient for generating an SSH private and public key. The output will return the private key, which should be kept somewhere private and inaccessible to others.

You can also upload a public key from a key pair using `openstack keypair create --public-key PATH_TO_PUBLIC_KEY KEY_NAME`, where `PATH_TO_PUBLIC_KEY` is the file path to the public key.

## Collect details about OpenStack environment

Show the role of each hardware node in an OpenStack cloud

```
$ openstack host list
```

Host Name	Service	Zone
hc1.edu.learnstack.org	scheduler	internal
hc2.edu.learnstack.org	scheduler	internal
hc3.edu.learnstack.org	scheduler	internal
hc1.edu.learnstack.org	conductor	internal
hc2.edu.learnstack.org	conductor	internal
hc3.edu.learnstack.org	conductor	internal
hc3.edu.learnstack.org	compute	nova
hc1.edu.learnstack.org	compute	nova
hc2.edu.learnstack.org	compute	nova

## Backup OpenStack Service and Ceph Configurations

This guide will explain how to make backups of OpenStack and Ceph configurations.

## Background Information

Our Private Cloud Core deployments make use of `kolla-ansible` and `ceph-ansible` to deploy the OpenStack services that comprise your cloud and the storage backend, Ceph.

OpenStack service configurations are kept within MariaDB databases as well as within the filesystem. An OpenStack service would be `nova`, `cinder`, or `glance`, to name a few examples.

## ***Backup OpenStack services using kolla-ansible***

Kolla-ansible allows you to back up the MariaDB databases associated with OpenStack services by using MariaDB's mariabackup function.

**NOTE!** - A good backup policy is to not store backups in the same location where production data lives, but rather to store backups offsite.

For an authoritative source of information on kolla-ansible and using mariabackup see these documentation links:

- [kolla-ansible](#)
- [MariaDB database backup and restore](#)

The [MariaDB database backup and restore guide](#) goes into full detail on how to enable backups, create them, and restore the backups.

---

## ***Using kolla-ansible***

To start using kolla-ansible, an environment needs to be created. This section explains the steps needed to create that environment.

---

To get a quick idea of what is required, here is a high-level overview of the steps:

```
# Copy kolla-ansible configuration from fm-deploy docker container
$ docker cp fm-deploy:/opt/kolla-ansible /opt/kolla-ansible

# Navigate to /opt/kolla-ansible
$ cd /opt/kolla-ansible

# Initialize a Python virtual environment
$ virtualenv .venv

# Activate the virtual environment
$ source .venv/bin/activate

# Install kolla-ansible using requirements.txt
$ pip install -r requirements.txt

# Set the SSH private key so kolla-ansible can connect to each host
$ export EXTRA_OPTS="--private-key /root/.ssh/fm-deploy"
```

The above takes care of preparing the kolla-ansible environment.

These steps are explained in more detail below.

---

## ***Prepare and use kolla-ansible***

Follow these steps to learn how to prepare and use kolla-ansible.

---

## ***Relevant files***

- Globals file: `/etc/kolla/globals.yml`
- Inventory file: `/etc/fm-deploy/kolla-ansible-inventory`

**Note!** – The above files may not be present on each host. Inspect each host until you find the above, and perform the kolla-ansible run from this host.

---

## Preparation

### Step 1 - Prepare environment

From the Docker container called `fm-deploy`, copy `/opt/kolla-ansible` to the local filesystem:

```
$ docker cp fm-deploy:/opt/kolla-ansible /opt/kolla-ansible
```

### Step 2 – Prepare Python virtual environment and install requirements

Create a Python virtual environment from which `kolla-ansible` will be used:

```
$ cd /opt/kolla-ansible
$ virtualenv .venv
$ source .venv/bin/activate
(.venv) $ pip install -r requirements.txt
```

### Step 3 – Export the private key

In order for `kolla-ansible` to make changes to each host, the SSH private key path needs to be set as an environment variable:

```
$ export EXTRA_OPTS="--private-key /root/.ssh/fm-deploy"
```

---

## Backup kolla-ansible Files

There are important files related to `kolla-ansible` that should be backed up.

These files are:

- `/etc/kolla/globals.yml`
- `/etc/kolla/passwords.yml`

Should they be lost, it will not be possible to manage the OpenStack cloud.

Create copies of these files and store them in an offsite backup location.

**NOTE!** – Due to the sensitive nature of these files, ensure the files are transported using an encrypted connection and are stored in a secure, private location.

---

## Create OpenStack Database Backups

With `kolla-ansible` prepared, you can create a full backup of all OpenStack service databases using `mariabackup`.

---

To create backups of the OpenStack databases, use:

```
$ kolla-ansible -i /etc/fm-deploy/kolla-ansible-inventory <command>
```

Note that <command> is a placeholder for the kolla-ansible command to run and in this case is mariadb\_backup.

---

The full command to create database backups is:

```
$ kolla-ansible -i /etc/fm-deploy/kolla-ansible-inventory mariadb_backup
```

This will create a Docker volume called **mariadb\_backup** that can be used to restore OpenStack databases.

The output of the kolla-ansible run will let you know to what host the backups were sent.

---

To see the Docker volume, SSH into the host where the backups were made, and run:

```
# docker volume ls | grep mariadb_backup
local          mariadb_backup
```

---

## Restore OpenStack Database Backups

This procedure outlines how to perform a full database restore. An incremental restoration can also be completed.

See [MariaDB's documentation](#) for information on how to use mariabackup to make incremental backups.

The previous kolla-ansible run created a Docker volume called **mariadb\_backup**. This volume can be used to restore the OpenStack service databases.

---

### Step 1 – Create Docker container

To restore this backup, create a new Docker container using the previously created volume:

```
# docker run --rm -it --volumes-from mariadb --name dbrestore \
--volume mariadb_backup:/backup kolla/centos-binary-mariadb:train-centos8 \
/bin/bash
```

### Step 2 – Perform full backup of databases

Once in that container, these series of commands can be run to perform a full backup of all OpenStack service databases:

```
$ cd /backup/
$ rm -rf /backup/restore/
$ mkdir -p /backup/restore/full
$ gunzip mysqlbackup-10-08-2020-1597091449.qp.xbc.xbs.gz
$ mbstream -x -C /backup/restore/full < mysqlbackup-10-08-2020-1597091449.qp.xbc.xbs
$ mariabackup --prepare --target-dir /backup/restore/full
```

### Step 3 – Stop MariaDB

Once the above commands are successfully run, stop the MariaDB Docker instance using:

```
# docker stop mariadb
```

**Step 4 – Restore database backups**

Navigate back into the container with the mariabackup volume mounted and either move or remove the contents of `/var/lib/mysql` and then move the contents of `/backup/restore/full` into `/var/lib/mysql`:

```
$ mkdir /backup/mariadb_original/
$ mv -v /var/lib/mysql/{*,\.[^\.]*} /backup/mariadb_original/
$ mv -v /backup/restore/full/{*,\.[^\.]*} /var/lib/mysql/
```

**Step 5 – Start MariaDB**

Finally, start the MariaDB Docker container:

```
# docker start mariadb
```

**OpenStack Configuration Backups**

In addition to backing up the MariaDB OpenStack databases, it is imperative the OpenStack service files are backed up. Services include nova, swift, and glance, to name a few.

All service configuration files are stored on the hardware nodes within the folder `/etc/kolla/`.

To backup the service's configuration files, copy `/etc/kolla` to an offsite backup location.

**Ceph Configuration Backups**

The configuration files for Ceph are stored on each hardware node within the folder `/etc/ceph`. Backing up this folder is sufficient for backing up Ceph's configuration.

To backup Ceph's configuration, copy `/etc/ceph` to an offsite backup location.

**Backing up Client Data and Disaster Recovery**

This guide will talk about ways client data can be backed up and how to restore that data in the event it becomes necessary.

It is suggested a good backup policy be in place in the event data needs to be restored. This policy will depend on various factors, such as importance of data, how much data needs to be stored, how far back in time to store data, what data to keep backups of, cost, and other factors that may be relevant to your OpenStack cloud.

Ceph acts as the distributed storage backend for this deployment of OpenStack. Ceph is naturally self-healing and there is no single point of failure, however it is still possible for it to fail, although this will be a rare occurrence. The more replicas used will only decrease the likelihood of failure at the expense of cost. This deployment of Ceph uses three replicas and is generally recommended as a good starting point.

If data loss is of the utmost importance, it is recommended that [RBD mirroring](#) within Ceph be setup. This means another Ceph cluster will have to exist and the data be mirrored between them. The additional ceph cluster can be setup in another data center to further decrease the failure domain.

## Where is client data currently stored?

All OpenStack client data is stored in Ceph pools.

This section will explain how to see the data stored in Ceph.

By running `rados lspools` from one of the OpenStack hardware nodes, you can see the individual ceph pools and what data is stored in each pool.

The following are the ceph pools where data is stored:

```
# rados lspools
device_health_metrics
images
volumes
vms
backups
metrics
manila_data
manila_metadata
.rgw.root
default.rgw.log
default.rgw.control
default.rgw.meta
default.rgw.buckets.index
```

The configuration for these pools is maintained using `ceph-ansible`.

You can see the contents of each pool by running `rbd ls -l POOL_NAME`.

An example listing the **images** pool:

```
# rbd ls -l images
NAME                                     SIZE      PARENT  FMT  PROT  LOCK
20b56b2d-1af5-46e1-a5c6-fa1f9a45245d    500 MiB
20b56b2d-1af5-46e1-a5c6-fa1f9a45245d@snap 500 MiB      2  yes
26a0fde5-69e7-4d85-ae4e-e167e295ecfa      0 B        2
26a0fde5-69e7-4d85-ae4e-e167e295ecfa@snap 0 B        2  yes
```

## How to create backups of volumes

### Using Horizon

Volumes can be backed up using Horizon. When backups are made they will be created within the **backups** ceph pool.

To create volume backups in Horizon, navigate to the **Volumes** tab, then to the next **Volumes** tab. This will display current volumes.

Find the dropdown next to the volume to back up, and click the “Create Backup” button.



Fill out the relevant details:

The volume backup should show as in progress:

Displaying 1 item

<input type="checkbox"/>	Name	Description	Size	Status	Volume Name	Snapshot	Actions
<input type="checkbox"/>	volume_1_backup	Backup for volume_1	20GB	Creating	40011aaa-3875-4236-9fd0-fff44c1fad21	-	

Displaying 1 item

And when done it should show listed like so:

Displaying 1 item

<input type="checkbox"/>	Name	Description	Size	Status	Volume Name	Snapshot	Actions
<input type="checkbox"/>	volume_1_backup	Backup for volume_1	20GB	Available	40011aaa-3875-4236-9fd0-fff44c1fad21	-	Restore Backup

Displaying 1 item

## Using OpenStackClient

Volume backups can also be created using OpenStackClient.

The following explains how to list volumes and make a volume backup.

List volumes and obtain UUID of volume to back up:

```
$ openstack volume list --fit-width
```

ID	Name	Status	Size	Attached
40011aaa-3875-4236-9fd0-fff44c1fad21	CentOS 8 (ce8-x86_64)	in-use	20	Attached
663419d7-df14-4472-9eb2-1f4c976103e9		in-use	20	Attached
e9d98e7b-5837-45bc-a0b1-5d9e50d7e686		in-use	20	Attached
057b53e7-eba9-4d2d-bec3-184240c59b29		available	20	
15848ac7-67db-460c-8be1-be1dcbb286f0		available	20	

Create volume backup of volume UUID 663419d7-df14-4472-9eb2-1f4c976103e9:

```
$ openstack volume backup create 663419d7-df14-4472-9eb2-1f4c976103e9 --force
```

Field	Value
id	0481a8ce-e571-45bc-b133-ac5496dce181
name	None

NOTE!: The `--force` flag is needed to make a backup of a volume that is in use.

List volume backups:



```
$ openstack volume backup list
```

ID	Name	Description	Status
0481a8ce-e571-45bc-b133-ac5496dce181	None	None	available

And finally for more detail on the volume backup created, you can use `openstack volume backup show UUID` where **UUID** is the volume backup UUID:

```
$ openstack volume backup show 0481a8ce-e571-45bc-b133-ac5496dce181
```

Field	Value
availability_zone	nova
container	backups
created_at	2020-12-15T16:55:39.000000
data_timestamp	2020-12-15T16:55:39.000000
description	None
fail_reason	None
has_dependent_backups	False
id	0481a8ce-e571-45bc-b133-ac5496dce181
is_incremental	False
name	None
object_count	0
size	20
snapshot_id	None
status	available
updated_at	2020-12-15T16:56:09.000000
volume_id	663419d7-df14-4472-9eb2-1f4c976103e9

### **Backing up ceph pool data**

Depending on your needs for data protection, it may be useful to back up data in the Ceph pools to a third party, such as an Amazon S3 bucket for example.

### **What to do in the event of a hard drive failure?**

There is no system in place that monitors for hardware failures, however it something being considered for future releases.

If it has been determined a hard drive has failed, a ticket will need to be created from the Flex Metal Central control panel. This ticket will be routed to our data center team who will replace the failed drive and will alert you when this task is done.

For now, monitoring of the hardware is something that will need to be set up. Multiple monitoring solutions exist already. [Icinga](#) or [Nagios](#) are two options that immediately stand out.

### **OpenStack Hardware Failures**

The purpose of this guide is to go over failure and maintenance scenarios in the OpenStack cloud that could occur and what to do to address them.

Things can and will go wrong. It is good to be prepared for these events.

### ***What should be done should a hardware node fail?***

If a hardware node fails or needs to come down for maintenance, you should know what steps to take. Depending on the maintenance required, you will either require the assistance of our data center staff or you can do the maintenance yourself.

### ***Planned maintenance***

This section describes the steps needed to take a hardware compute node out of the cloud in the event work needs to be done on it or the cloud needs to be reduced in size.

**NOTE!** - If you know a node or nodes need maintenance that require a hardware modification you'll need to create a ticket from the Flex Metal Central control panel to our data center staff to perform that task for you.

For official documentation on this subject, see OpenStack's [Compute Node Failures and Maintenance](#) guide.

The general work flow for bringing a compute node down will involve first disabling that node, finding the instances on that node, migrating those instances to another node, and removing any ceph Object Storage Daemons (OSDs). Optionally, you can migrate the instances back to the original node when the maintenance is done.

OpenStackClient will be required to perform the maintenance.

### ***Procedure for removing a compute node***

Start with disabling the `nova-compute` service on the appropriate node:

```
$ openstack compute service set --disable --disable-reason \
maintenance COMPUTE_NODE_NAME nova-compute
```

List the instances on that node:

```
$ openstack server list --host COMPUTE_NODE_NAME --all-projects
```

Migrate the instances to another node:

```
$ openstack server migrate INSTANCE_UUID --live-migration
```

**NOTE** - This deployment of OpenStack is using ceph as the backend shared storage so there is no need to pass the `--block-migration` flag to `openstack server migrate`.

Because OpenStack has been deployed using `kolla-ansible`, each OpenStack service runs in a docker container.

Stop the `nova_compute` docker container:

```
# docker stop nova_compute
```

Perform the needed maintenance, and then restart the `nova_compute` service:

```
# docker start nova_compute
```

Verify the `nova_compute` docker container is running:

```
# docker ps | grep nova_compute
286e1b2e2ae5      kolla/centos-binary-nova-compute:train-centos8
"dumb-init --single-..." 2 months ago      Up 18 minutes
nova_compute
```

Finally, verify the nova service has connected to the messaging service, AMQP:

```
$ grep AMQP /var/log/kolla/nova/nova-compute.log
```

## Unplanned maintenance

There are times where unplanned maintenance is required. This section will describe what can be done in the event a compute node goes down unexpectedly.

The primary concern is that instances associated with the compute node that has failed will no longer work.

---

## Ceph failure scenarios and recovery

Ceph by nature is resilient to hardware failure and self-healing.

The primary concern with ceph is failed hard drives. How can an operator be alerted to a failed hard drive? Will ceph continue to function if a drive is lost?

Generally, ceph will continue to function if a drive is lost, however the drive should be replaced as soon as possible.

## How do you know if a hard drive has failed?

Currently there is no monitoring for failed ceph drives, however the intention is to monitor for these events in the future. Due to this, it is recommended monitoring of drives be put into place. Software such as [Icinga](#) or [Nagios](#) are viable options for monitoring.

If it is suspected a drive has failed, you should first determine if this really is the case.

The overall procedure for determining if a drive has failed is to:

- Check Ceph health
- See if the OSD associated with the drive in question can be started if it is stopped
- Check the OSD's mount point using `df -h`
- Use `smartctl` on the drive in question

The following explains these steps in more detail.

---

### Procedure

#### Reference:

[https://access.redhat.com/documentation/en-us/red\\_hat\\_ceph\\_storage/4/html/operations\\_guide/handling-a-disk-failure](https://access.redhat.com/documentation/en-us/red_hat_ceph_storage/4/html/operations_guide/handling-a-disk-failure)

From one of the hardware nodes, perform the following checks:

Check if ceph is healthy:

```
# ceph health
```

Find the location of the OSD within the CRUSH map:

```
# ceph osd tree | grep -i down
```

On the node that houses the OSD, try to start the OSD using `systemctl` where **OSD** is a placeholder for the actual OSD identifier:

```
# systemctl start ceph-osd@OSD.service
```

The systemctl unit file for the OSD will vary depending on which OSD has failed. In this case the systemctl unit file is called `ceph-osd@0.service`.

If a hard drive has failed, our data center team will need to replace it. A ticket will need to be made in Flex Metal Central to alert the team of the failure. The drive or drives will be replaced by our team.

## Adding and Removing OpenStack Hardware Nodes

If you find the cloud growing, it may come time to add another node to the cloud. Similarly, you may not need hardware any longer.

This guide will demonstrate how to add or remove a node from an OpenStack cloud. This is only applicable to those who have added an additional **Storage and Compute** node via Flex Metal Central. If you have not added any additional nodes that means each node is a control plane node and those nodes should **not** be removed.

## Types of Hardware Nodes

The following are the current nodes provided through Flex Metal Central.

- **Cloud Core - Standard**
- **Storage and Compute - Standard**

Here is a screenshot showing three **Cloud Core** nodes and one **Storage and Compute** node:

Hardware						
Class	Type	Storage	Cores	RAM	Hostname	
Cloud Core - Standard	mb_standard_v1	3 TB	16	128 GB	eager-sarah.local	
Cloud Core - Standard	mb_standard_v1	3 TB	16	128 GB	busy-josephb.local	
Cloud Core - Standard	mb_standard_v1	3 TB	16	128 GB	pensive-michaelc.local	
Storage and Compute - Standard	mb_standard_v1	3 TB	16	128 GB	perfect-lobster.local	

Note that the three **Cloud Core** nodes are part of the original cloud. The **Storage and Compute** node was added to the cloud.


## Adding a node

Hardware nodes can be added to an existing cloud using Flex Metal Central.

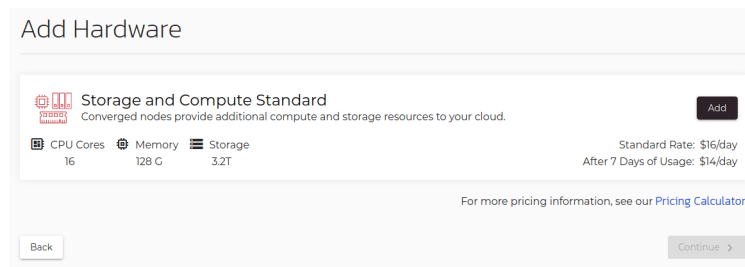
There is currently only one node type that can be added which is **Storage and Compute - Standard**. Future updates will see new node types.

Once logged into Flex Metal Central, navigate to the cloud you are working with. Find the button near the top right that says “Add Hardware”. Clicking this button will allow you to add additional hardware to the cloud.

**Add Hardware:**



**Add a new Storage and Compute - Standard node:**



## Removing a Storage and Compute Node

There is no automated way to safely remove a compute node. In order to safely remove it you will need to ensure all data is removed from the node, including instances, Ceph data, and anything else.

**NOTE!** – When you are ready to remove the node, you can do so by clicking the three vertical dots, then the **Remove** link, next to the node in Flex Metal Central. You will see a screen asking you to confirm removal of the node which you must confirm before doing so. It is strongly advised you are sure everything is removed from the node before doing this.

This section will demonstrate the steps needed to remove a hardware node. Note that this is a manual process.

**NOTE!** – If a compute node needs to be removed and you do not feel comfortable migrating data off of it, it is recommended you reach out to support who can help perform this task for you.

There are a number of things that need to be checked and considered before removing a hardware node. In addition the type of hardware node will dictate what needs to occur. This section will describe the general process to remove a **Storage and Compute** node.

## Requirements before removing a node

- This task can only be performed using the command line.
- Instances running on the node will need to be migrated to another compute node.
- Ceph OSDs need to be removed from ceph.
- Ceph ansible needs its inventory updated
- OpenStack needs to be updated that a node will be removed.
- Kolla-ansible will need its inventory file updated to reflect removal of the node

The following will go into detail on how to perform each step.

## Remove a Storage and Compute node

With the above in mind, this section will expand upon the steps in the previous section.

**Reference:** <https://docs.openstack.org/kolla-ansible/latest/user/adding-and-removing-hosts.html>

**Procedure****Live migrate instances:**

You will want to first find all instances on the node, confirm the receiving node has enough resources to host the instances, then perform a live migration of those instances.

**NOTE!** – Since Ceph is used for disk storage you do not need to account for disk space when moving instances from one compute node to another. You will only need to account for RAM and VCPUs.

**Step 1 – Disable instance scheduling on node**

The node being removed will need to have instance scheduling disabled. This causes it so no new instances can be created on this node.

To perform this, run `openstack compute service set HOST nova-compute --disable`.

For example, to disable instance scheduling for the `perfect-lobster.local` host run:

```
$ openstack compute service set perfect-lobster.local nova-compute --disable
```

**Step 2 – Collect instance information**

This guide will assume you have a compute node called `perfect-lobster.local` that needs to be removed from the OpenStack cloud.

In addition, there are three instances on this node:

- `migrate_me-1`
- `migrate_me-2`
- `migrate_me-3`

You can use `$ openstack server list` with additional flags to obtain the details of the instances on `perfect-lobster.local`.

For example:

```
$ openstack server list --host perfect-lobster.local -f value -c ID -c Name
a248c6b2-c4f5-4b5a-82ca-0dc71edd9757 migrate_me-2
ad489cd4-4c0c-42c7-aeae-21da6c00693b migrate_me-3
d143fb86-c7a4-4bc8-b6dc-7969097ef34b migrate_me-1
```

The above makes use of **value** output formatting and specifies the **ID** and **Name** columns as output.

**Step 3 – Determine compute host**

With the instance information acquired, the next step is to determine what compute host these instances can be migrated to.

To obtain all compute hosts, you can use `$ openstack compute service list`.

For example:

```
$ openstack compute service list
```

ID	Binary	Host	Zone	Status	State	Updated At
12	nova-scheduler	eager-sarahl.local	internal	enabled	up	2021-01-28T23
51	nova-scheduler	busy-josephb.local	internal	enabled	up	2021-01-28T23
66	nova-scheduler	pensive-michaelcu.local	internal	enabled	up	2021-01-28T23
3	nova-conductor	eager-sarahl.local	internal	enabled	up	2021-01-28T23
12	nova-conductor	busy-josephb.local	internal	enabled	up	2021-01-28T23
27	nova-conductor	pensive-michaelcu.local	internal	enabled	up	2021-01-28T23
30	nova-compute	eager-sarahl.local	nova	enabled	up	2021-01-28T23

33	nova-compute	busy-josephb.local	nova	enabled	up	2021-01-28T23
36	nova-compute	pensive-michaelcu.local	nova	enabled	up	2021-01-28T23
37	nova-compute	perfect-lobster.local	nova	enabled	up	2021-01-28T23

What is shown in the above output is there are four total compute hosts. One of them (`perfect-lobster.local`) is being removed from the cluster.

This example will select `eager-sarah1.local` as the new host to migrate the instances to.

#### Step 4 – Check available host resources

You must first ensure this host has enough resources to contain these instances. To do so, you can run `$ openstack host show HOSTNAME`.

For example:

```
$ openstack host show eager-sarah1.local
```

Host	Project	CPU	Memory MB	Disk GB
eager-sarah1.local	(total)	16	121026	11923
eager-sarah1.local	(used_now)	4	6144	25
eager-sarah1.local	(used_max)	2	2048	20
eager-sarah1.local	b9e8639372014c0b85cbfaffa6e1b5a8	2	2048	20

#### Step 5 – Find instance resource usage

To ensure the instances will fit you will also need to know what resources they consume. This means you will need to know the flavor set for each instance.

You can run something like this to get the flavor for each instance on the `perfect-lobster.local` node:

```
$ openstack server list --host perfect-lobster.local -f value -c ID -c Name -c Flavor
a248c6b2-c4f5-4b5a-82ca-0dc71edd9757 migrate_me-2 hcl1.micro
ad489cd4-4c0c-42c7-aeae-21da6c00693b migrate_me-3 hcl1.micro
d143fb86-c7a4-4bc8-b6dc-7969097ef34b migrate_me-1 hcl1.micro
```

This shows the `hcl1.micro` flavor is used by each instance on this host.

To know what resources are consumed by this flavor use `$ openstack flavor show FLAVOR`:

```
$ openstack flavor show hcl1.micro
```

Field	Value
OS-FLV-DISABLED:disabled	False
OS-FLV-EXT-DATA:ephemeral	0
access_project_ids	None
disk	10
id	hcl1.micro
name	hcl1.micro
os-flavor-access:is_public	True
properties	
ram	1024
rxtx_factor	1.0
swap	1024
vcpus	1

This output reveals the amount of VCPUs, RAM, and disk space allocated.

To get tidier output of the flavor details you can run:

```
$ openstack flavor show hcl.micro -c disk -c ram -c vcpus
+-----+-----+
| Field | Value |
+-----+-----+
| disk  | 10    |
| ram   | 1024  |
| vcpus | 1     |
+-----+-----+
```

From this information, it can be determined the receiving host will require 3 VCPUs, 3GB of RAM, and 30GB of disk space since there are three instances to migrate and each is using the `hcl.micro` flavor. Remember, however, that disk space does not need to be accounted for since ceph is used for data storage and the data storage is shared across each node.

## Step 6 – Live migrate the instances

You can now safely migrate these instances to the `eager-sarah1.local` node.

The base command to perform the live migration is:

```
$ openstack --os-compute-api-version 2.79 server migrate \
--live-migration --host HOSTNAME INSTANCE_UUID
```

The command to live migrate these three instances for this demonstration is:

```
openstack server list --host perfect-lobster.local -f value -c ID | while read id; do
  openstack --os-compute-api-version 2.79 server migrate --live-migration --host eager-sarah1.local $id
done
```

You can confirm the status of the live migration by using `$ openstack server list`:

```
$ openstack server list
+-----+-----+-----+-----+
| ID                                     | Name                               | Status      | Networks |
+-----+-----+-----+-----+
| a248c6b2-c4f5-4b5a-82ca-0dc71edd9757 | migrate_me-2                       | ACTIVE      | Internal= |
| ad489cd4-4c0c-42c7-aeae-21da6c00693b | migrate_me-3                       | MIGRATING   | Internal= |
| d143fb86-c7a4-4bc8-b6dc-7969097ef34b | migrate_me-1                       | MIGRATING   | Internal= |
| 85033a0f-66c6-41d4-b679-c7350da2685f | openstackclient_js_demo            | ACTIVE      | Internal= |
| 0052cd0f-70fb-4cf7-8b13-2bec350c0e51 | openstackclient_jumpstation        | ACTIVE      | Internal= |
+-----+-----+-----+-----+
```

Here it can be seen two of the instances are being migrated.

## Step 7 – Confirm live migration success

To confirm the live migration completed successfully and the instances are on the new host, you can use something similar to:

```
openstack server list -f value -c ID -c Name | grep migrate_me | while read id name; do
  echo "$name $(openstack server show -f value -c 'OS-EXT-SRV-ATTR:host' $id)"
done
```

The following is the actual output from running the above command:

```
migrate_me-2 eager-sarah1.local
migrate_me-3 eager-sarah1.local
migrate_me-1 eager-sarah1.local
```

This indicates the live migration was successful.

Confirm no instances remain on the original host:

```
$ openstack server list --host perfect-lobster.local
```



If the above returns no output, the live migration was a complete success and you can move on to the next step.

---

## Remove Ceph OSDs:

After the instances have been migrated, it is time to remove the node's OSDs from the Ceph cluster.

**Reference:** <https://docs.ceph.com/en/latest/rados/operations/add-or-rm-osds/>

This continues to assume the compute host being removed is `perfect-lobster.local`.

You will need to determine what OSDs are on this host and remove them from ceph.

### Step 1 – Determine OSDs

From any of the hardware nodes you can use `# ceph osd tree` to find which OSD is on a particular host.

Example:

```
# ceph osd tree
```

ID	CLASS	WEIGHT	TYPE	NAME	STATUS	REWEIGHT	PRI-AFF
-1		11.64398	root	default			
-5		2.91100	host	busy-josephb			
1	ssd	2.91100		osd.1	up	1.00000	1.00000
-3		2.91100	host	eager-sarah1			
0	ssd	2.91100		osd.0	up	1.00000	1.00000
-7		2.91100	host	pensive-michaelcu			
2	ssd	2.91100		osd.2	up	1.00000	1.00000
-9		2.91100	host	perfect-lobster			
3	ssd	2.91100		osd.3	up	1.00000	1.00000

This indicates the host `perfect-lobster.local` has only one OSD, with **ID** of 3. This is the OSD that will need to be removed.

### Step 2 – Remove the OSD

To remove the OSD, use the command `ceph osd out OSD_NUMBER`.

For example:

```
# ceph osd out 3
marked out osd.3.
```

Following that, watch ceph's status using `ceph -w` to ensure the cluster returns to a healthy state (output truncated):

```
# ceph -w
cluster:
  id:      a08c2963-1d75-42ef-be50-1fc61419623d
  health:  HEALTH_WARN
           Degraded data redundancy: 441/13998 objects degraded (3.150%), 17 pgs degraded

[...]

io:
  recovery: 68 MiB/s, 14 keys/s, 11 objects/s

progress:
  Rebalancing after osd.3 marked out (1s)
  [.....]

2021-01-29T15:43:20.780381+0000 mon.eager-sarah1 [WRN] Health check update: Degraded data re
```

```
2021-01-29T15:43:20.817317+0000 mon.eager-sarah1 [INF] Health check cleared: PG_DEGRADED (wa
2021-01-29T15:43:20.817342+0000 mon.eager-sarah1 [INF] Cluster is now healthy
```

The above shows the cluster in a degraded state as the OSD is outed and that Ceph returns to a healthy state. This indicates no issues occurred with outing the OSD.

### Step 3 – Stop the OSD

With the OSD removed, the OSD `systemctl` service on the `perfect-lobster.local` host needs to be stopped.

In this case, the unit file for this OSD is called `ceph-osd@3.service`.

Stop the service:

```
# systemctl stop ceph-osd@3.service
```

### Step 4 – Remove OSD from ceph configuration

The OSD will now need to be removed from ceph's configuration.

To do so, you will need to use `ceph osd purge OSD_NUMBER --yes-i-really-mean-it`:

```
# ceph osd purge 3 --yes-i-really-mean-it
purged osd.3
```

### Step 5 – Remove OSD from Ceph crush map

The OSD will still need to be removed from the crush map.

To remove the `perfect-lobster` OSD from the crush map, use:

```
# ceph osd crush rm perfect-lobster
```

### Step 6 – Confirm OSD has been removed

You can use `ceph osd tree` to confirm the OSD from the node has been removed.

For example:

```
# ceph osd tree
ID CLASS WEIGHT TYPE NAME STATUS REWEIGHT PRI-AFF
-1  8.73299 root default
-5  2.91100 host busy-josephb
 1  ssd 2.91100 osd.1 up 1.00000 1.00000
-3  2.91100 host eager-sarah1
 0  ssd 2.91100 osd.0 up 1.00000 1.00000
-7  2.91100 host pensive-michaelcu
 2  ssd 2.91100 osd.2 up 1.00000 1.00000
```

Here it can be seen the `perfect-lobster.local` host is no longer present in this output.

These steps take care of the Ceph changes that are required.

---

### Clean up services on remaining nodes:

The next steps are to clean up the host from the compute and network service listing within OpenStack.

The following demonstrate removing the host `perfect-lobster.local` from the network agent list and the compute service list.

Remove host from network agent list:

```
openstack network agent list --host perfect-lobster.local -f value -c ID | while read id; do
  openstack network agent delete ${id}
done
```

Remove host from compute service list:

```
openstack compute service list --os-compute-api-version 2.53 --host perfect-lobster.local -f
openstack compute service delete --os-compute-api-version 2.53 ${id}
done
```

Confirm the `perfect-lobster.local` has been removed from both section by using:

```
$ openstack network agent list --host perfect-lobster.local
$ openstack compute service list --os-compute-api-version 2.53 --host \
perfect-lobster.local
```

---

### Stop the services running on the node (stop the Docker service):

The next step will be to stop the Docker service on the node being removed.

Use the following to stop the Docker service:

```
# systemctl status docker
```

---

### Submit the request in Flex Metal Central to remove the node:

At this point, you are ready to request the node removal in Flex Metal Central.

Login to Flex Metal Central and go the **Manage** section for the cloud. From here you will see a listing of the hardware nodes. Select the node being removed and find the three vertical dots to the right of it. Clicking this will bring up the option to remove the node. Click **Remove** to initiate that process.

**NOTE!** – At this time, when you request a node to be removed, it will create a ticket and our support team will handle that request manually and follow up when that is done.

---

### Test the remaining nodes to ensure everything still functions:

Once at this point, the node removal is complete. At this time, it is recommended you test general functionality of the cloud, such as can instances still be created, do the hardware nodes still respond to ping, and the like.

## Automating OpenStack

It is possible to automate deploying various OpenStack services, such as instance creation. Terraform is software that can be used to do this.

See InMotion Hosting's [documentation](#) on using Terraform with OpenStack through Flex Metal Cloud.

## Miscellaneous Guides

This guide stores miscellaneous information and may be useful, however the information in this guide is not needed to go through setting up the cloud.

---

## Associate an instance with a private network

This section will explain how to take an instance created on the provider network and associate that with a private network. This may not be a very common task assuming an instance was created on the appropriate network to begin with.

In the event an instance was created but on the wrong network, you can have it associated with another network.

In Horizon, pull up the listing on instances under the **Compute** heading on the left, then find **Instances** under that. Find the instance you're working with then from the dropdown on the far right, choose the **Attach Interface** option.

### Attach Interface:

<input type="checkbox"/>	Instance Name	Image Name	IP Address	Flavor	Key Pair	Status	Availability Zone	Task	Power State	Age	Actions	
<input type="checkbox"/>	Server 1	CentOS 8 (c88-x86_64)	173.231.202.94	hc1.small	ssh_key	Active	us-east-1	nova	None	Running	6 days, 19 hours	<div>Create Snapshot</div> <div>Associate Floating IP</div> <div>Attach Interface</div> <div>Detach Interface</div>

Displaying 1 item

Under the **Network** dropdown choose the appropriate network. This example will choose the **192.168.0.0/24** network.

### Attach Interface form:

Attach Interface

The way to specify an interface \*

by Network (and IP address)

Network\*

Internal (192.168.0.0/24)

Fixed IP Address ⓘ

Description:

Select the network for interface attaching.

Cancel

Attach Interface

With that done, the instance is now associated with two networks. To finish moving the instance to the new network, use the same drop down in the instance listing page to detach an interface. Locate the option called **Detach Interface**, then choose the network from the drop down and submit the form to remove it.